

Big Data Management

HBase **(BigTable)**



HBase: Overview

- **HBase is a distributed column-oriented data store built on HDFS**
- **HBase is Apache open source project for storage for Hadoop (BigTable Clone)**
- **Data is logically organized into tables, rows and columns**

HBase

- A distributed data store that can scale horizontally to 1,000s of commodity servers and petabytes of indexed storage.
- Designed to operate on top of Hadoop distributed file system (HDFS) for scalability, fault tolerance, and high availability.

HBase vs. HDFS

- Both distributed systems scale to thousands of nodes
- **HDFS** is good for batch processing (scans over big files)
 - Not good for record lookup
 - Not good for incremental addition of small batches or updates
- **HBase** is designed to efficiently address the above
 - Fast record lookup
 - Support for record-level insertion & updates (but not in place)
- **HBase updates done by creating new versions of values**

HBase Is Not ...

- Tables have one primary index, the *row key*.
- There are three types of lookups (select columns):
 - Fast lookup using row key and optional timestamp.
 - Full table scan
 - Range scan from region start to end.
 - (**No join operators !**)
- Limited atomicity and transaction support:
 - HBase supports multiple batched mutations of single rows only.
- Not accessed or manipulated via SQL:
 - Programmatic access via Java, REST, or Thrift APIs.

HBase vs. HDFS

	Plain HDFS/MR	HBase
Write pattern	Append-only	Random write, bulk incremental
Read pattern	Full table scan, partition table scan	Random read, small range scan, or table scan
Hive (SQL) performance	Very good	4-5x slower
Structured storage	Do-it-yourself / TSV / SequenceFile / Avro / ?	Sparse column-family data model
Max data size	30+ PB	~1PB

If application has neither random reads/writes → Stick to HDFS

HBase vs. RDBMS

	RDBMS	HBase
Data layout	Row-oriented	Column-family-oriented
Transactions	Multi-row ACID	Single row only
Query language	SQL	get/put/scan/etc *
Security	Authentication/Authorization	Work in progress
Indexes	On arbitrary columns	Row-key only
Max data size	TBs	~1PB
Read/write throughput limits	1000s queries/second	Millions of queries/second

HBase Data Model

(Google's BigTable Model)

Tables are sorted by Row

Table schema define its *column families* .

- Each family consists of any number of columns

- Each column consists of any number of versions

- Columns only exist when inserted, NULLs are free.

- Columns within family are sorted & stored together

Everything except table names are byte[]

(Row, Family: Column, Timestamp) → Value

HBase: Keys & Column Families

Each row has a Key

Each record is divided into Column Families

The diagram illustrates the structure of the PERSON TABLE in HBase. It is organized into columns representing different data categories. The first column is the 'row key', which contains 'PersonID' values ranging from 1 to 500,000,000. The subsequent columns represent 'column families': 'personal_data' (containing 'Name' and 'Address'), 'demographic' (containing 'BirthDate' and 'Gender'), and an ellipsis indicating more families. Arrows from the text above point to these components: one to the row key column, and two to the column family headers. A specific cell in the 'Address' column (Merlin, Stonehenge, England) is also highlighted with an arrow.

row key	personal_data		demographic		...
PersonID	Name	Address	BirthDate	Gender	...
1	H. Houdini	Budapest, Hungary	1926-10-31	M	
2	D. Copper	New Jersey, USA	1956-09-16	M	
3	Merlin	Stonehenge, England	1136-12-03	F	
...	
500,000,000	F. Cadillac	Nevada, USA	1964-01-07	M	

Figure 2 - Census Data in Column Families

Each column family consists of one or more Columns

HBase Physical Model

- Each column family is stored in a separate file (called *HTables*)
- Key & Version numbers are replicated with each column family
- Empty cells are not stored

HBase maintains a multi-level index on values:
<key, column family, column name, timestamp>

Table 5.3. ColumnFamily contents

Row Key	Time Stamp	ColumnFamily "contents:"
"com.cnn.www"	t6	contents:html = "<html>..."
"com.cnn.www"	t5	contents:html = "<html>..."
"com.cnn.www"	t3	contents:html = "<html>..."

Table 5.2. ColumnFamily anchor

Row Key	Time Stamp	Column Family anchor
"com.cnn.www"	t9	anchor:cnnsi.com = "CNN"
"com.cnn.www"	t8	anchor:my.look.ca = "CNN.com"

Column Families

- Different sets of columns may have different properties and access patterns
- Configurable by column family:
 - Compression (none, gzip, LZO)
 - Version retention policies
 - Cache priority
- CFs stored separately on disk: access one without wasting IO on the other.

HBase Regions

- Each HTable (column family) is partitioned horizontally into *regions*
 - Regions are counterpart to HDFS blocks

Table 5.3. ColumnFamily contents

Row Key	Time Stamp	ColumnFamily "contents:"
"com.cnn.www"	t6	contents:html = "<html>..."
"com.cnn.www"	t5	contents:html = "<html>..."
"com.cnn.www"	t3	contents:html = "<html>..."

⋮



Each will be one region

When to use Hbase?

- You need random write, random read, or both (*but not neither*)
- You need to do many thousands of operations per second on multiple TB of data
- Your access patterns are well-known and simple