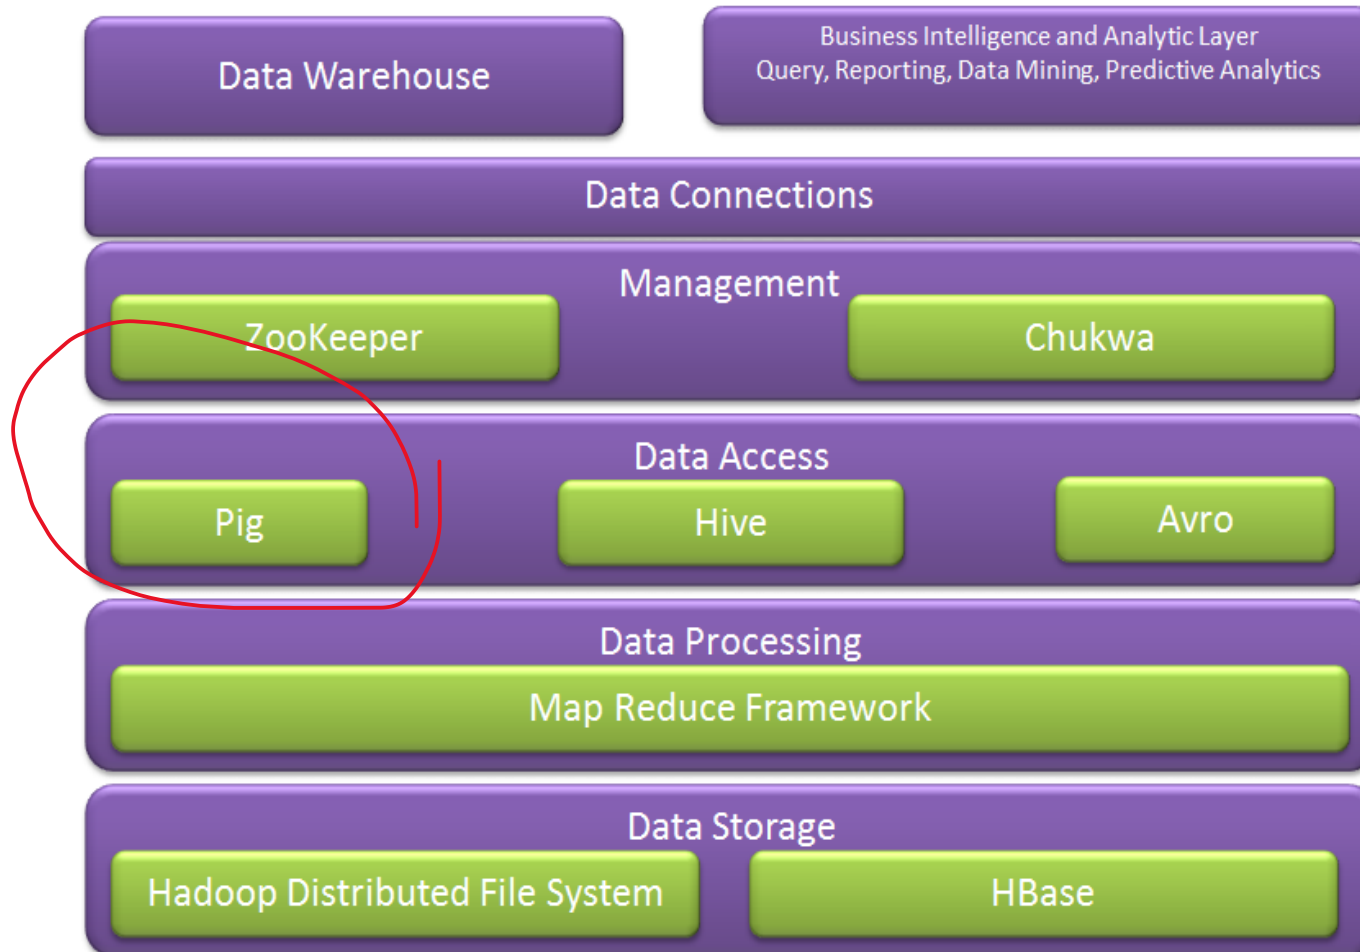


MapReduce High-Level Languages:

PIG

Hadoop Ecosystem



Query Languages for Hadoop

- **Java:** Hadoop's Native Language
- **Pig:** Query and Workflow Language (Yahoo)
- **Hive:** SQL-Based Language (Facebook)
- **HBase:** Column-oriented DB for MapReduce

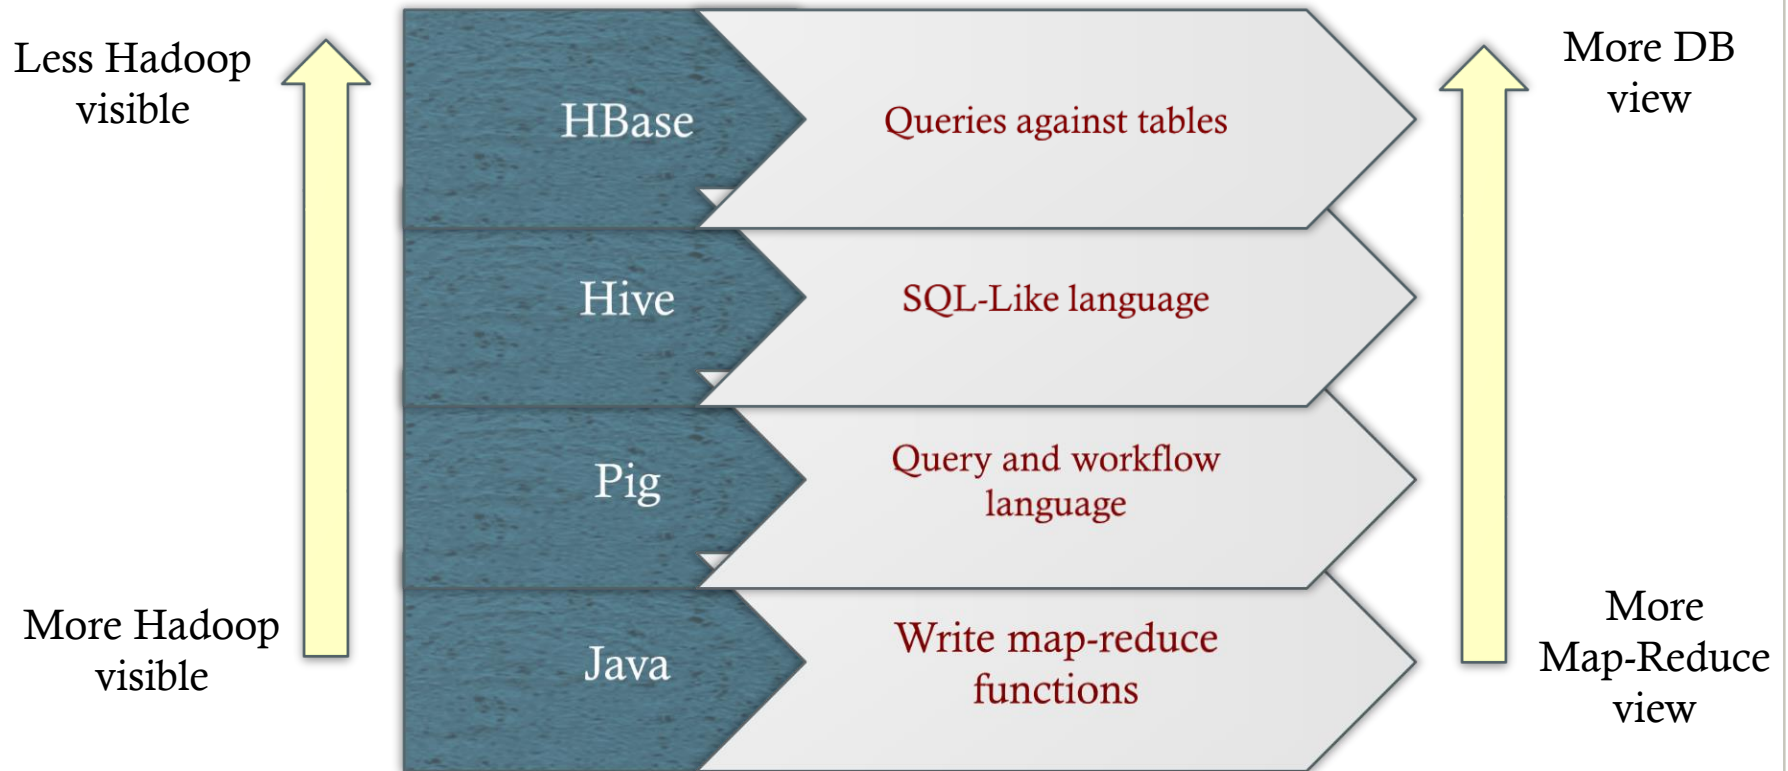
Java is Hadoop's Native Language

- Hadoop itself is written in Java

 **FULL CONTROL**

- **Provided Java APIs**
 - For mappers, reducers, combiners, partitioners
 - Input and output formats
- **Other languages, e.g., Pig or Hive, convert their queries to Java MapReduce code**

Levels of Abstraction



Apache Pig

(Chapter 16, Hadoop: The Definitive Guide)

What is Apache Pig ?



- A platform for analyzing large data sets with a **high-level language for expressing** data analysis programs.
- Compiles down to MapReduce jobs
- Open-source language
- Developed by Yahoo

High-Level Language

```
raw = LOAD 'excite.log' USING PigStorage('\t') AS (user, id, time, query);
clean1 = FILTER raw BY id > 20 AND id < 100;
clean2 = FOREACH clean1 GENERATE
    user, time,
    org.apache.pig.tutorial.sanitize(query) as query;
user_groups = GROUP clean2 BY (user, query);
user_query_counts = FOREACH user_groups
    GENERATE group, COUNT(clean2), MIN(clean2.time), MAX(clean2.time);
STORE user_query_counts INTO 'uq_counts.csv' USING PigStorage(',');
```


Pig Components

Two Main Components

- High-level language (Pig Latin)
 - Set of commands
- Two execution modes
 - Local: reads/write to local file system
 - Mapreduce: connects to Hadoop cluster and reads/writes to HDFS

Two Modes

- Interactive mode
 - Console
- Batch mode
 - Submit a script

Why Language like Pig?

- Common design patterns as KEY WORDS
 - joins, distinct, counts
- Data flow analysis
 - A script can map to multiple map-reduce jobs
- Avoids Java-level errors
 - not everyone can write java code
- Can be interactive mode
 - Issue commands and get results

Example I: More Details

Read file from HDFS

The input format (text, tab delimited)

Define run-time schema

```
raw = LOAD 'excite.log' USING PigStorage('\t') AS (user, id, time, query);  
clean1 = FILTER raw BY id > 20 AND id < 100;  
clean2 = FOREACH clean1 GENERATE  
        user, time,  
        org.apache.pig.tutorial.sanitize(query) as query;  
user_groups = GROUP clean2 BY (user, query);  
user_query_counts = FOREACH user_groups  
    GENERATE group, COUNT(clean2), MIN(clean2.time), MAX(clean2.time);  
STORE user_query_counts INTO 'uq_counts.csv' USING PigStorage(',');
```

Filter the rows on predicates

For each row, do some transformation

Grouping of records

Compute aggregation for each group

Store the output in a file

Text, Comma delimited

Pig: Language Features

- **Keywords**
 - Load, Filter, Foreach Generate, Group By, Store, Join, Distinct, Order By
- **Aggregations**
 - Count, Avg, Sum, Max, Min
- **Schema**
 - Defined at query-time (not when files are loaded)
- **Extension of Logic**
 - UDFs
- **Data**
 - Packages for common input/output formats

Example2: Parameterized Template

Script can take arguments

Data are “ctrl-A” delimited

Define types of the columns

```
A = load '$widerow' using PigStorage('\u0001')  
      as (name: chararray, c0: int, c1: int, c2: int);
```

```
B = group A by name parallel 10;
```

Specify the need of 10 reduce tasks

```
C = foreach B generate group, SUM(A.c0) as c0, SUM(A.c1) as c1,  
      AVG(A.c2) as c2;
```

```
D = filter C by c0 > 100 and c1 > 100 and c2 > 100;
```

```
store D into '$out';
```

Example 3: Partition Join

Register UDFs & custom inputformats

```
register pigperf.jar;
```

Function the jar file to read the input file

```
A = load 'page_views' using org.apache.pig.test.udf.storefunc.PigPerformanceLoader()  
as (user, action, timespent, query_term, timestamp, estimated_revenue);
```

```
B = foreach A generate user, (double) estimated_revenue;
```

Load the second file

```
alpha = load 'users' using PigStorage('\u0001') as (name, phone, address, city, state, zip);
```

```
beta = foreach alpha generate name, city;
```

Join the two datasets (40 reducers)

```
C = join beta by name, B by user parallel 40;
```

```
D = group C by $0;
```

Group after the join (can reference columns by position)

```
E = foreach D generate
```

```
store E into 'L3out'
```

This join and grouping,
how many map-reduce jobs ?

Example 3: Partition Join

```
register pigperf.jar;
```

```
A = load 'page_views' using org.apache.pig.test.udf.storefunc.PigPerformanceLoader()  
    as (user, action, timespent, query_term, timestamp, estimated_revenue);
```

```
B = foreach A generate user, (double) estimated_revenue;
```

```
alpha = load 'users' using PigStorage('\u0001') as (name, phone, address, city, state, zip);
```

```
beta = foreach alpha generate name, city;
```

```
C = join beta by name, B by user parallel 40;
```

```
D = group C by $0;
```

```
E = fore
```

```
store E i
```

This grouping can be done in the same map-reduce job because it is on the same key
(Pig can do this optimization !)

Example 4: Replicated Join

```
register pigperf.jar;
```

```
A = load 'page_views' using org.apache.pig.test.udf.storefunc.PigPerformanceLoader()  
    as (user, action, timespent, query_term, timestamp, estimated_revenue);
```

```
Big = foreach A generate user, (double) estimated_revenue;
```

```
alpha = load 'users' using PigStorage('\u0001') as (name, phone, address, city, state, zip);
```

```
small = foreach alpha generate name, city;
```

```
C = join Big by user, Small by name using 'replicated'
```

```
store C into 'out';
```

Map-only join
(small dataset is the second)

Optimization in joining a big
dataset with a small one

Example 5: Multiple Outputs

```
A = LOAD 'data' AS (f1:int,f2:int,f3:int);
```

```
DUMP A;
```

```
(1,2,3)
```

```
(4,5,6)
```

```
(7,8,9)
```

Split the records into sets

```
SPLIT A INTO X IF f1<7, Y IF f2==5, Z IF (f3<6 OR f3>6);
```

```
DUMP X;
```

```
(1,2,3)
```

```
(4,5,6)
```

Dump command to display the data

```
DUMP Y;
```

```
(4,5,6)
```

Store multiple outputs

```
STORE X INTO 'x_out';
```

```
STORE Y INTO 'y_out';
```

```
STORE Z INTO 'z_out';
```

Run independent jobs in parallel

```
D1 = load 'data1' ...
```

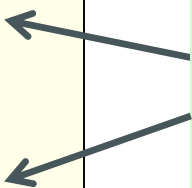
```
D2 = load 'data2' ...
```

```
D3 = load 'data3' ...
```

```
C1 = join D1 by a, D2 by b
```

```
C2 = join D1 by c, D3 by d
```

C1 and C2 are
two independent
jobs that can run
in parallel



Pig Latin vs. SQL

- Pig Latin is procedural (dataflow programming model)
 - Step-by-step query style is easier to write for some
- SQL is declarative but not step-by-step style

SQL

```
insert into ValuableClicksPerDMA
select dma, count(*)
from geoinfo join (
    select name, ipaddr
    from users join clicks on (users.name = clicks.user)
    where value > 0;
) using ipaddr
group by dma;
```

Pig
Latin

```
Users          = load 'users' as (name, age, ipaddr);
Clicks         = load 'clicks' as (user, url, value);
ValuableClicks = filter Clicks by value > 0;
UserClicks     = join Users by name, ValuableClicks by user;
Geoinfo        = load 'geoinfo' as (ipaddr, dma);
UserGeo        = join UserClicks by ipaddr, Geoinfo by ipaddr;
ByDMA          = group UserGeo by dma;
ValuableClicksPerDMA = foreach ByDMA generate group, COUNT(UserGeo);
store ValuableClicksPerDMA into 'ValuableClicksPerDMA';
```

Pig Latin vs. SQL

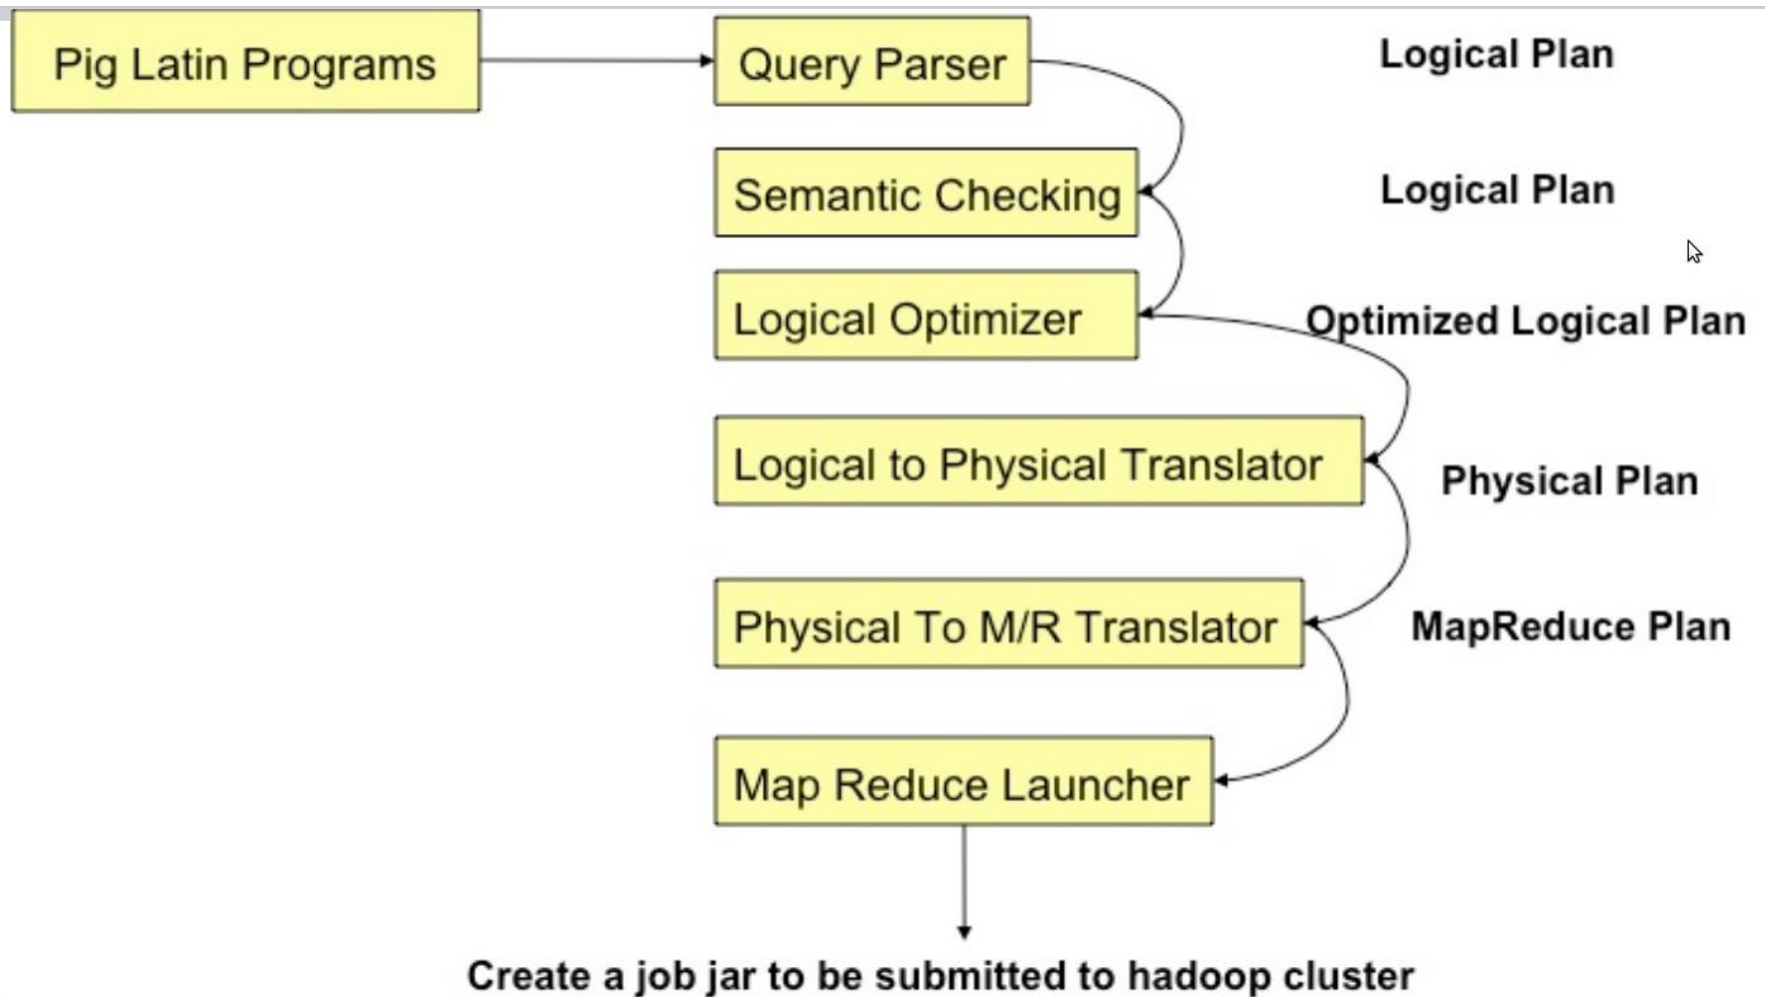
- **In Pig Latin**

- Lazy evaluation (data not processed prior to STORE command)
- Data can be stored at any point during the pipeline
- Schema and data types are lazily defined at run-time
- An execution plan can be explicitly defined by users (via hints)
 - Use optimizer hints (due to the lack of complex optimizers)

- **In SQL:**


- Query plans are solely decided by the system (powerful opt)
- Data cannot be stored in the middle (or, at least not user-accessible)
- Schema and data types are defined at the creation time

Pig Compilation



Logic Plan

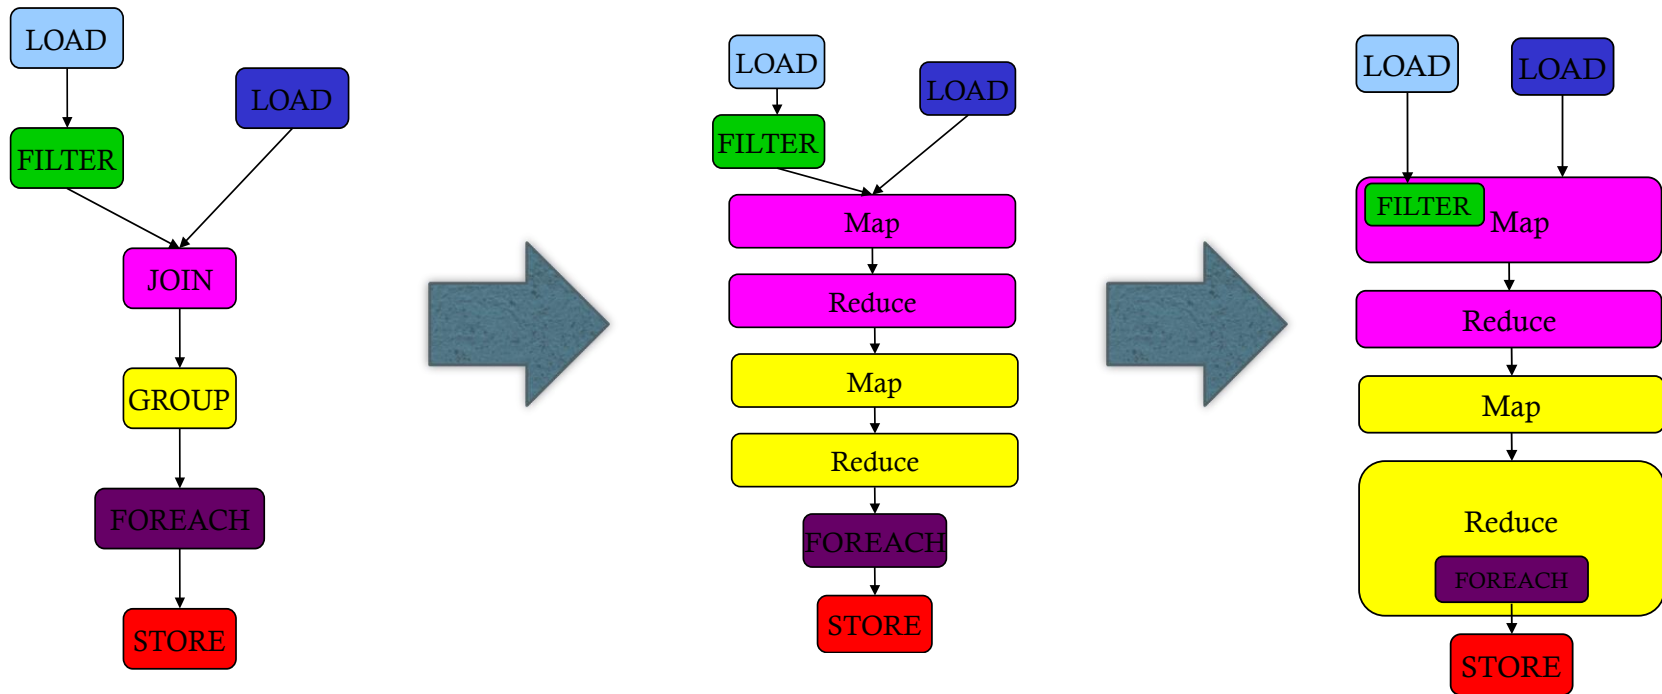
```
A=LOAD 'file1' AS (x, y, z);  
B=LOAD 'file2' AS (t, u, v);  
C=FILTER A by y > 0;  
D=JOIN C BY x, B BY u;  
E=GROUP D BY z;  
F=FOREACH E GENERATE  
  group, COUNT(D);  
STORE F INTO 'output';
```



Physical Plan

- Mostly 1:1 correspondence with logical plan
- **Except for:**
 - Join, Distinct, (Co)Group, Order
- Some optimizations are done automatically


Generation of Physical Plans




If the Join and Group By are on the same key →
The two map-reduce jobs would be merged into one.

Java vs. Pig

1/20 the lines of code



1/16 the development time



Performance is comparable
(Java is slightly better)

Pig References

- **Pig Tutorial**
 - <http://pig.apache.org/docs/r0.7.0/tutorial.html>
- **Pig Latin Reference Manual 2**
 - http://pig.apache.org/docs/r0.7.0/piglatin_ref1.html
- **Pig Latin Reference Manual 2**
 - http://pig.apache.org/docs/r0.7.0/piglatin_ref2.html
- **PigMix Queries**
 - <https://hpccsystems.com/why-hpcc-systems/benchmarks/pigmix-hpcc>

Apache Pig

