


Hadoop/MapReduce: Overall Computing Paradigm

Part 1

Large-Scale Data Analytics

- MapReduce computing paradigm (E.g., Hadoop) vs. Traditional database systems




vs.



- **Enterprises are turning to Hadoop**
 - Especially applications generating *big data*
 - Web applications, social networks, scientific applications

Why Hadoop is able to compete?



vs.



Scalability (petabytes of data, thousands of machines)



Flexibility in accepting all data formats (no schema)



Efficient and simple fault-tolerant mechanism



Commodity inexpensive hardware



Performance (tons of indexing, tuning, data organization tech.)



Fast Processing

Interactive processing



Transactions and consistency guarantees (ACID)

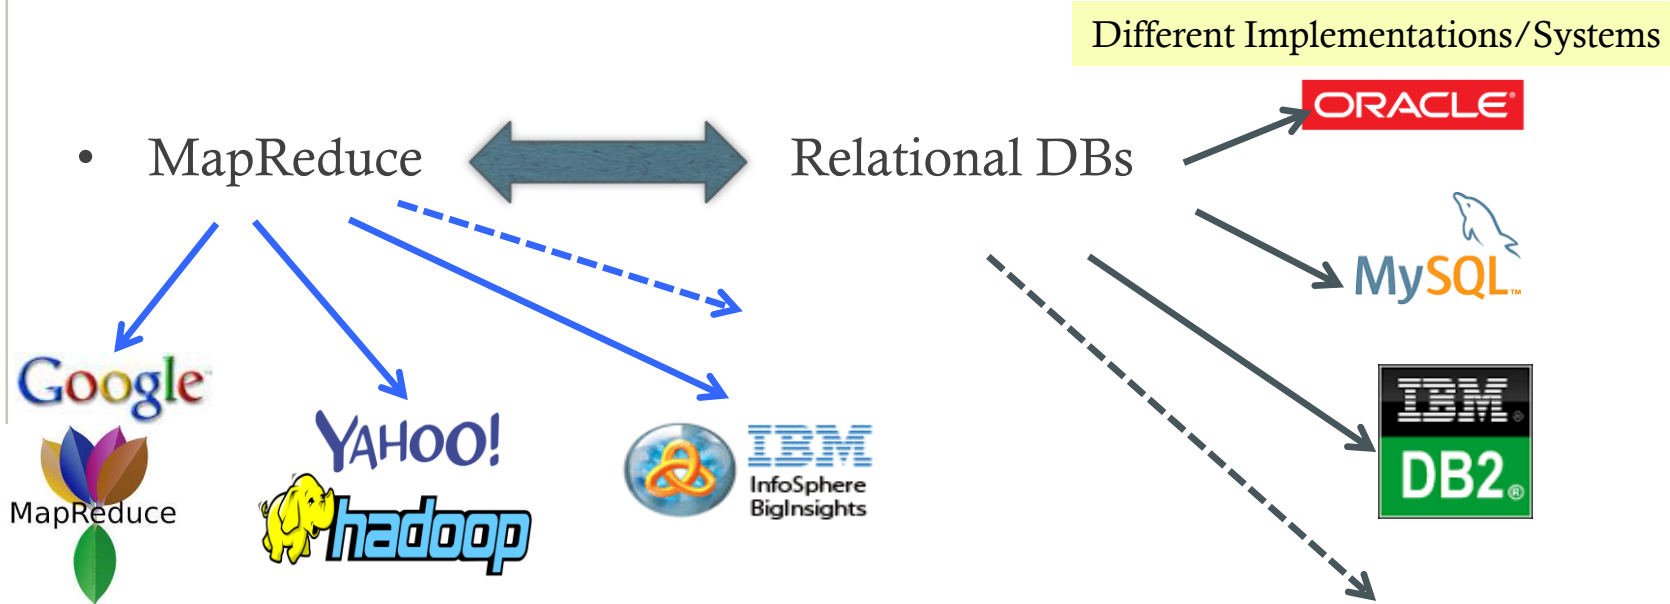


Features:

- Provenance tracking
- Annotation management
-

What is MapReduce

- **MapReduce is a “computational paradigm”**
 - A specific mechanism of processing the data



What is Hadoop ?

- Hadoop is a software framework for *distributed processing* of *large datasets* across *large clusters* of computers
 - *Large datasets* → Terabytes or petabytes of data
 - *Large clusters* → hundreds or thousands of nodes
- Hadoop is open-source implementation for Google **MapReduce**
- Hadoop based on simple programming model called *MapReduce*
- Hadoop is based on a simple data model, *any data will fit*

Compute Cluster

A rack of N machines




One machine \longleftrightarrow One node


Compute Cluster

- Cluster → Set of machines connected together


Shared-memory



Shared-disk




Shared-nothing

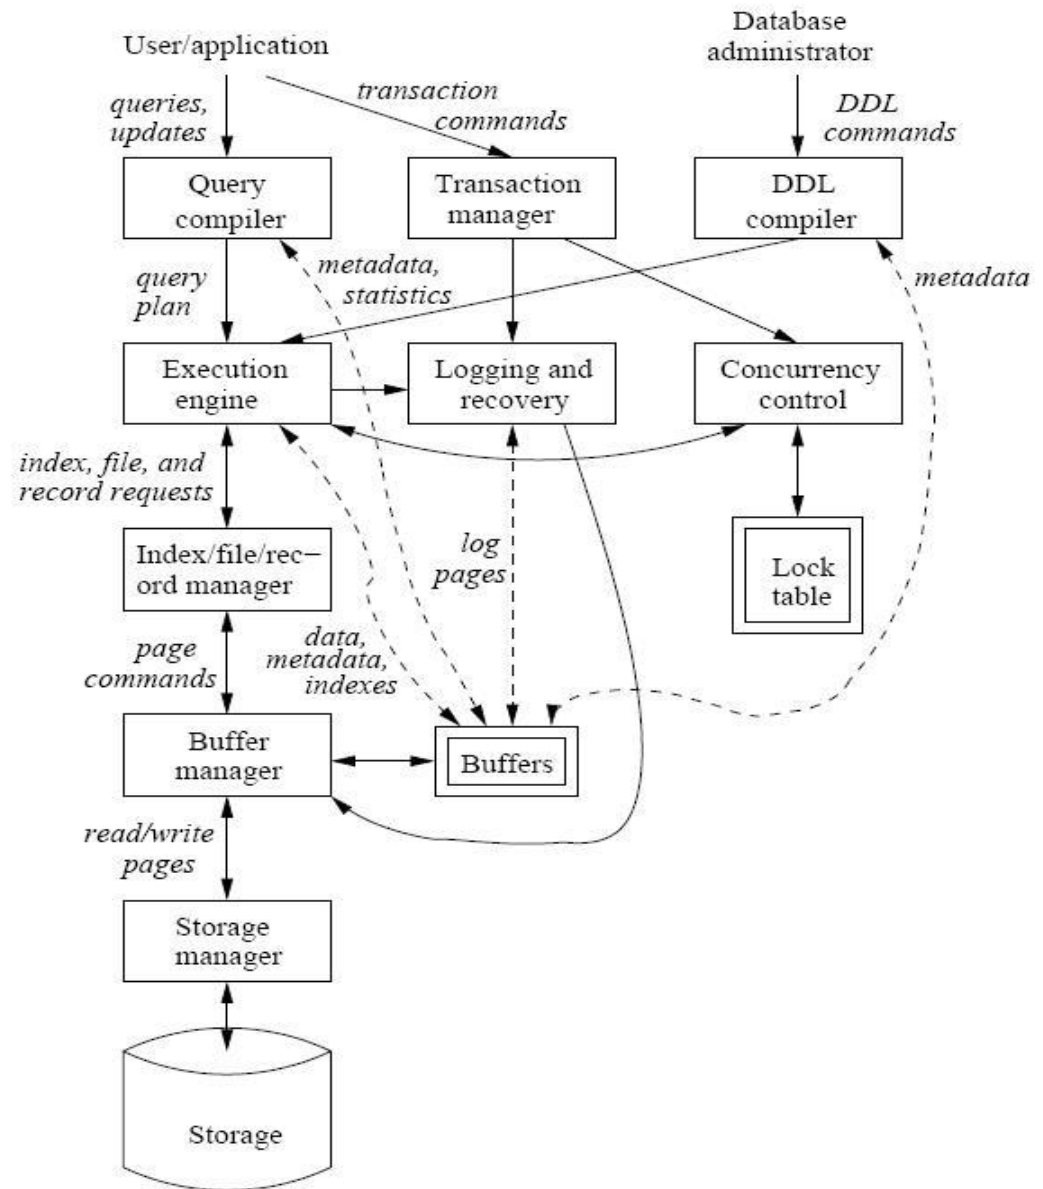


What is Hadoop (Cont'd)

- **Hadoop framework consists of two main layers:**
 - Distributed file system (HDFS)
 - Execution engine (MapReduce)



Contrast it with RDBMS



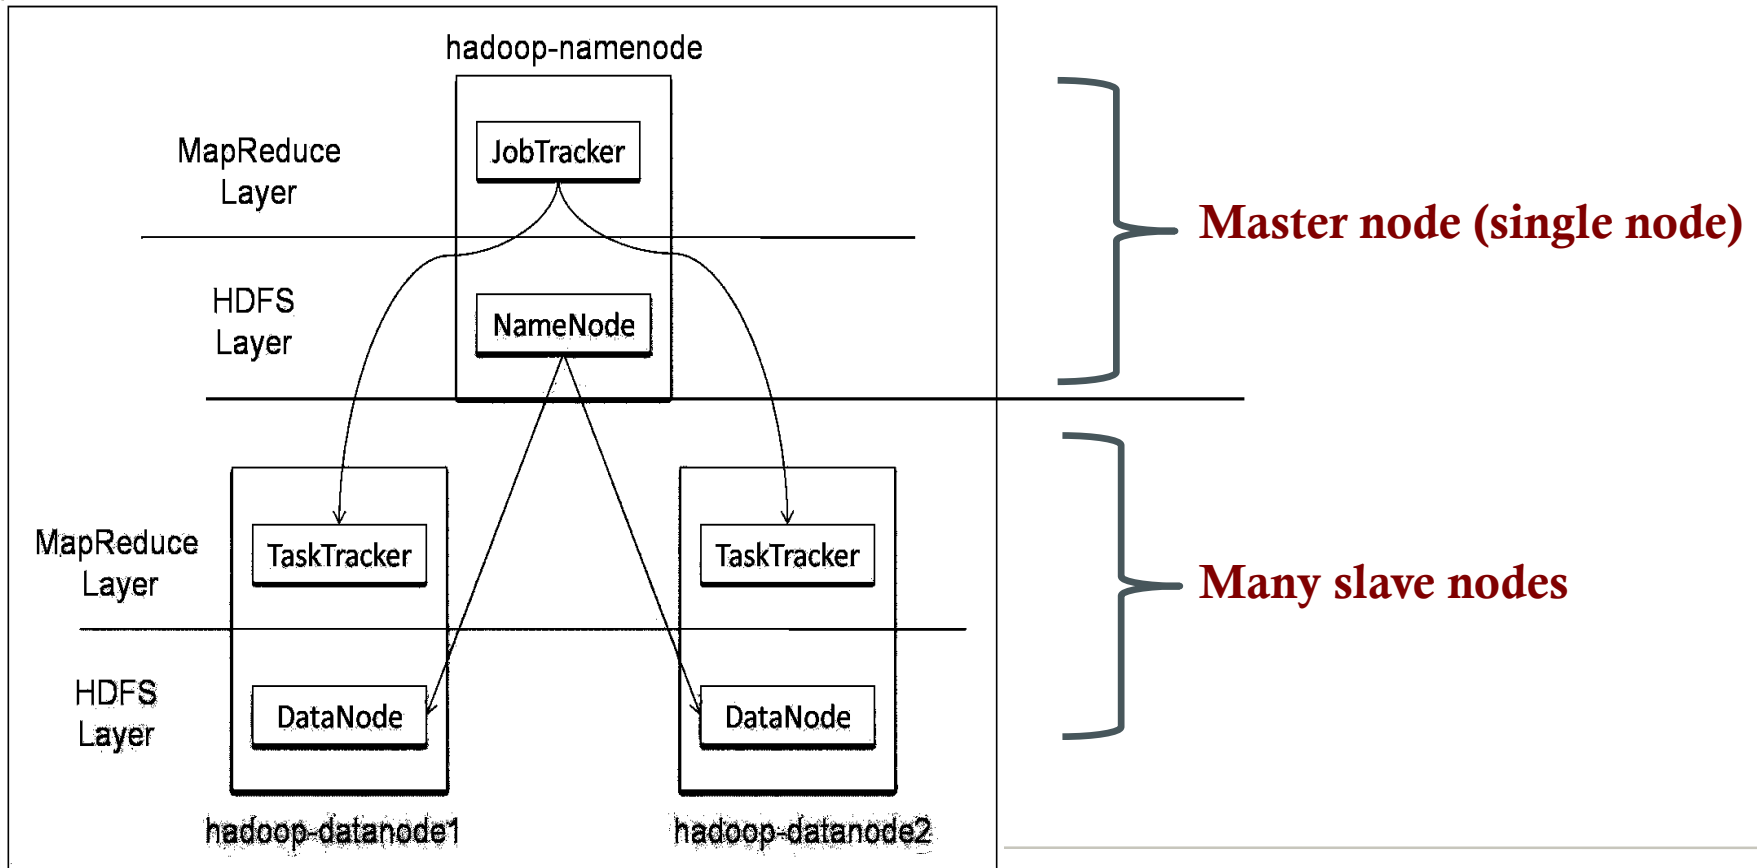
Database management system components

Design Principles of Hadoop

- **Scale to big data**
- **Simple to use**
- **Minimal functionality**

Hadoop Master/Slave Architecture

- Hadoop is designed as *master-slave shared-nothing* architecture



Design Choices of Hadoop

- Need to process big data
- Need to parallelize computation across thousands of nodes
- **Commodity hardware**
 - Large number of low-end cheap machines working in parallel to solve a computing problem
- This is in contrast to **Parallel DBs**
 - Small number of high-end expensive machines

Design Principles of Hadoop

- **Automatic parallelization & distribution**
 - Hidden from the end-user
- **Fault tolerance and automatic recovery**
 - Nodes/tasks will fail and will recover automatically
- **Clean and simple programming abstraction**
 - Users only provide two functions “map” and “reduce”

Who Uses MapReduce/Hadoop?

- Google: Inventors of MapReduce computing paradigm
- Yahoo: Developing Hadoop open-source of MapReduce
- IBM, Microsoft, Oracle
- Facebook, Amazon, AOL, NetFlix
- Many others ++
- Universities and research labs