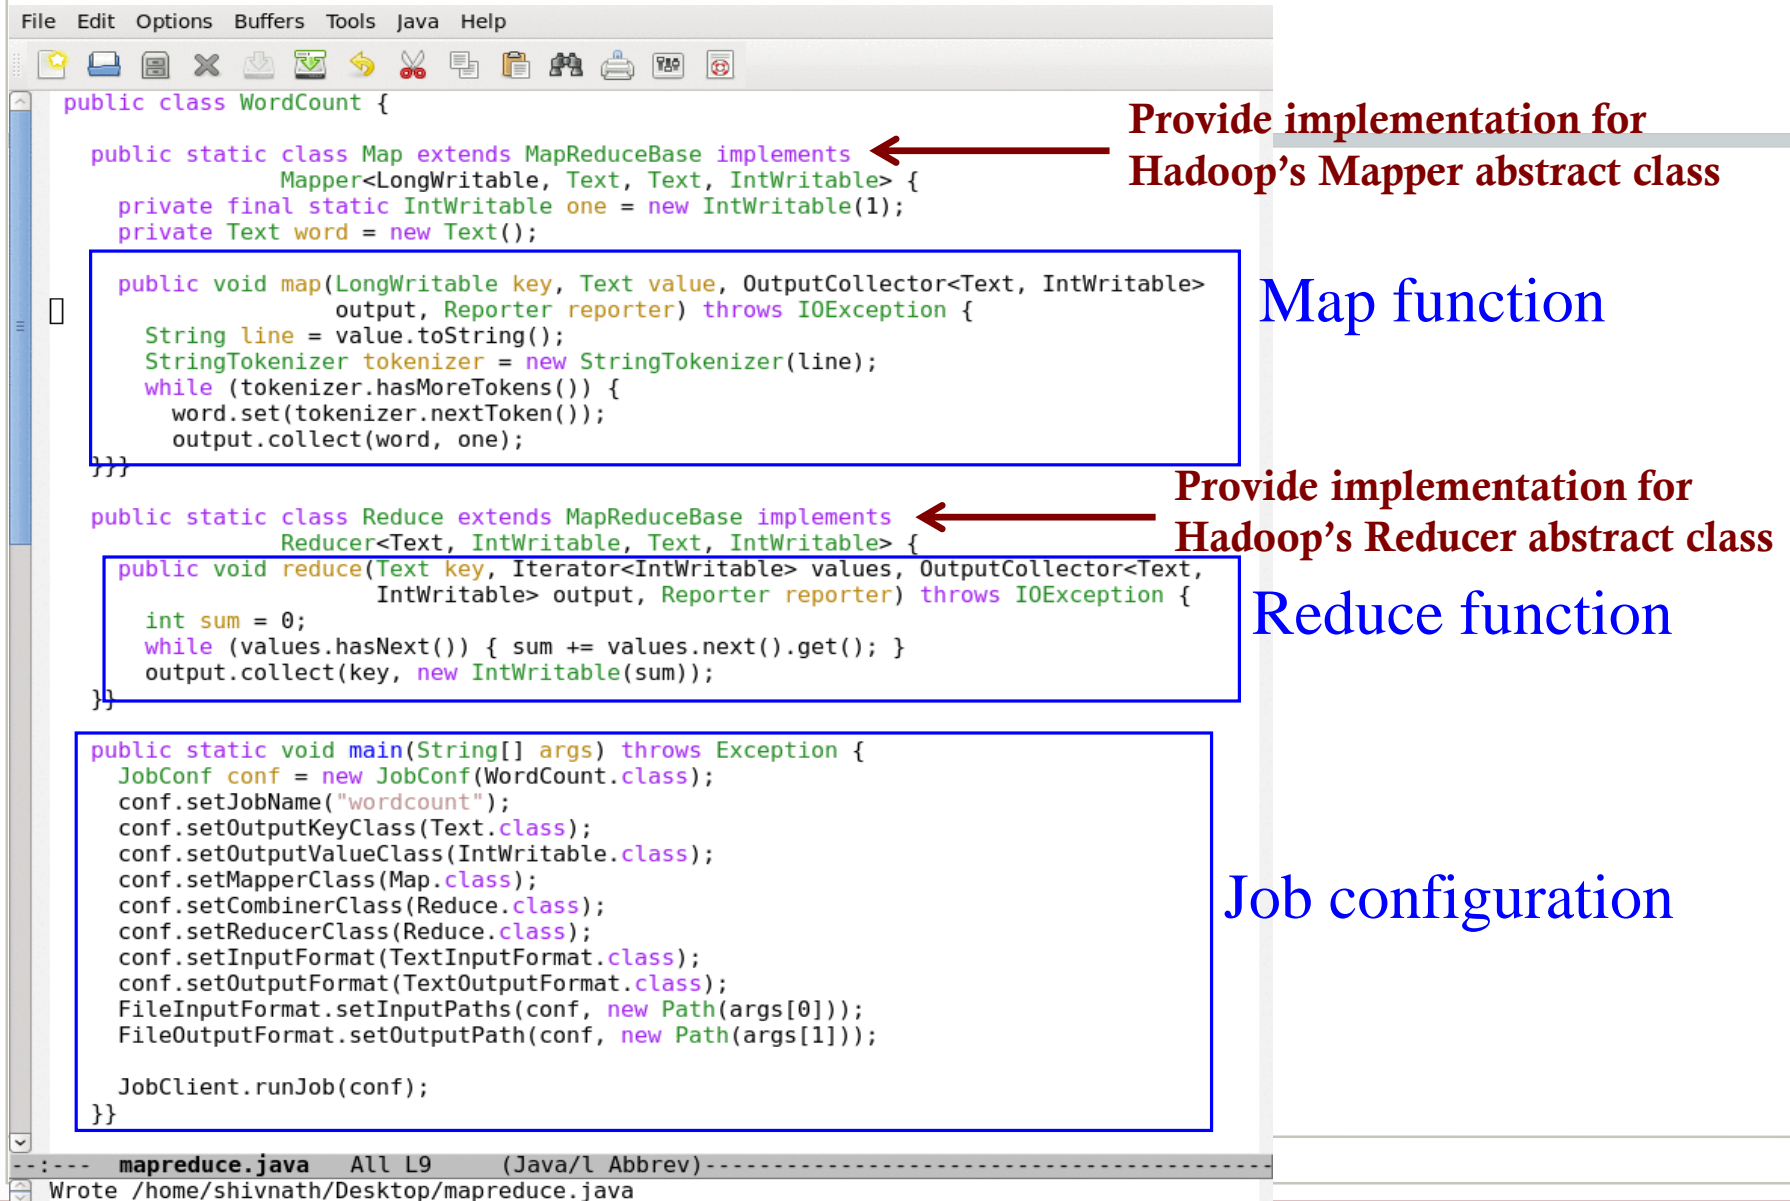


# Hadoop/MapReduce: Optimizations

# How it looks like in Java



The image shows a screenshot of a Java IDE with the file `mapreduce.java` open. The code is for a Hadoop MapReduce application named `WordCount`. Annotations with arrows point to specific parts of the code:

- Provide implementation for Hadoop's Mapper abstract class**: Points to the `Map` class definition: `public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {`
- Map function**: Points to the `map` method implementation: `public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {`
- Provide implementation for Hadoop's Reducer abstract class**: Points to the `Reduce` class definition: `public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {`
- Reduce function**: Points to the `reduce` method implementation: `public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {`
- Job configuration**: Points to the `main` method where the Hadoop job is configured and run: `public static void main(String[] args) throws Exception {`

```
public class WordCount {  
  
    public static class Map extends MapReduceBase implements  
        Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>  
            output, Reporter reporter) throws IOException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                output.collect(word, one);  
            }  
        }  
    }  
  
    public static class Reduce extends MapReduceBase implements  
        Reducer<Text, IntWritable, Text, IntWritable> {  
        public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,  
            IntWritable> output, Reporter reporter) throws IOException {  
            int sum = 0;  
            while (values.hasNext()) { sum += values.next().get(); }  
            output.collect(key, new IntWritable(sum));  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        JobConf conf = new JobConf(WordCount.class);  
        conf.setJobName("wordcount");  
        conf.setOutputKeyClass(Text.class);  
        conf.setOutputValueClass(IntWritable.class);  
        conf.setMapperClass(Map.class);  
        conf.setCombinerClass(Reduce.class);  
        conf.setReducerClass(Reduce.class);  
        conf.setInputFormat(TextInputFormat.class);  
        conf.setOutputFormat(TextOutputFormat.class);  
        FileInputFormat.setInputPaths(conf, new Path(args[0]));  
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
  
        JobClient.runJob(conf);  
    }  
}
```

---:--- mapreduce.java All L9 (Java/L Abbrev) ---:---  
Wrote /home/shivnath/Desktop/mapreduce.java

# Optimizations

---

- **Four optimizations for map-reduce processing**

# Optimization 1

- In Color Count example, what if number of colors is small  
→ then can we optimize the map-side ?



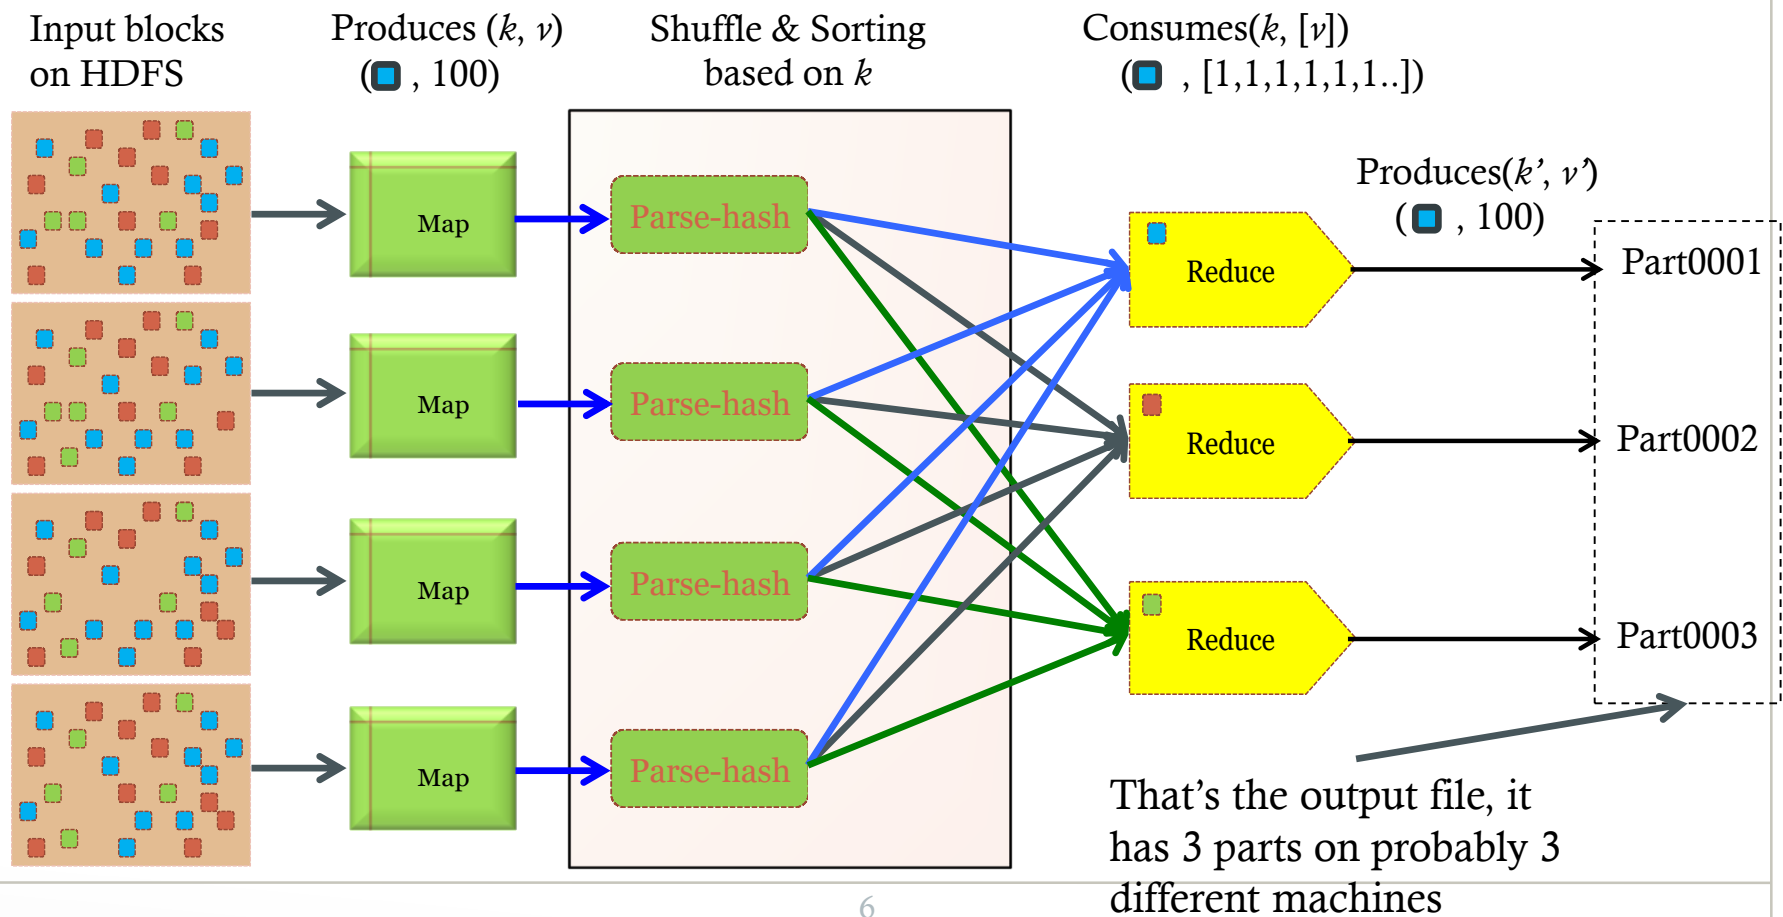
- Each map function can have a small main-memory hash table (color, count)
- With each line, update the hash table and produce nothing
- When done, report each color and its local count

	10
	5
	7
	20

**Gain:** Reduce amount of shuffled/sorted data over network

# Optimization 1: Takes Place inside Mappers

**Saves network messages (Typically very expensive phase)**



# Optimization 1

- Small main-memory hash table (color, count)
- Update hash table
- Report result using hash table when done.

	10
	5
	7
	20

**Q1: Where to build the hash table?**

**Q2: How to know when done?**

**Q3: What do, when done?**

# Inside the Mapper Class

Called once after all records done  
(Here you can produce the output)

## Method Summary

protected void	<code>cleanup(<a href="#">Mapper.Context</a> context)</code> Called once at the end of the task.
-------------------	---

protected void	<code>map(<a href="#">KEYIN</a> key, <a href="#">VALUEIN</a> value, <a href="#">Mapper.Context</a> context)</code> Called once for each key/value pair in the input split.
-------------------	---

Called for each record

void	<code>run(<a href="#">Mapper.Context</a> context)</code> Expert users can override this method for more complete control over the execution of the Mapper.
------	---


protected void	<code>setup(<a href="#">Mapper.Context</a> context)</code> Called once at the beginning of the task.
-------------------	---

Called once before any record processed  
(Here you can build the hash table)

Reducer has similar functions...

# Opt. 2: Map-Combine-Reduce

- On each machine, we partially aggregate results from mappers

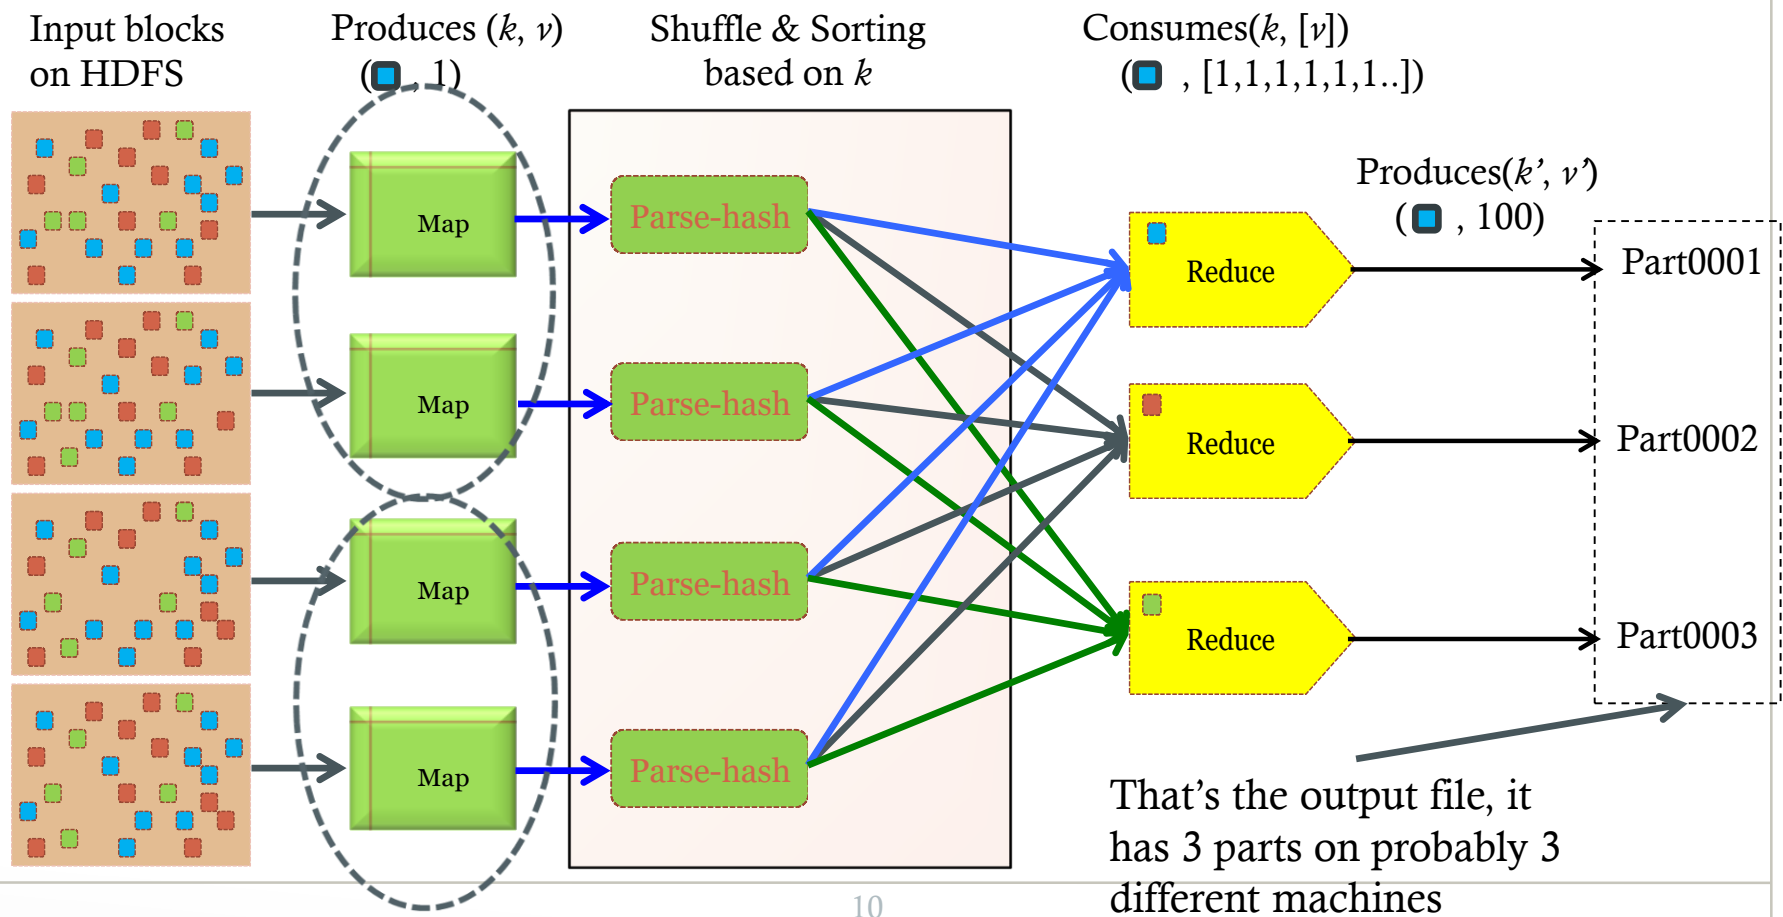


- A **combiner** is a **reducer** that runs on each machine to locally aggregate (*via user code*) mappers' outputs from this machine
- Combiners' output is shuffled and sorted for 'real' reducers



# Optimization 2: Outside Mappers, But on Each Machine

**Combiner runs on each node to partially aggregate local mappers' output**



# Tell Hadoop to use a Combiner

```
File Edit Options Buffers Tools Java Help
public class WordCount {
    public static class Map extends MapReduceBase implements
        Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>
            output, Reporter reporter) throws IOException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                output.collect(word, one);
            }
        }
    }

    public static class Reduce extends MapReduceBase implements
        Reducer<Text, IntWritable, Text, IntWritable> {
        public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,
            IntWritable> output, Reporter reporter) throws IOException {
            int sum = 0;
            while (values.hasNext()) { sum += values.next().get(); }
            output.collect(key, new IntWritable(sum));
        }
    }

    public static void main(String[] args) throws Exception {
        JobConf conf = new JobConf(WordCount.class);
        conf.setJobName("wordcount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(Map.class);
        conf.setCombinerClass(Reduce.class);
        conf.setReducerClass(Reduce.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));

        JobClient.runJob(conf);
    }
}
```

mapreduce.java All L9 (Java/L Abbrev)

Wrote /home/shivnath/Desktop/mapreduce.java

Not all jobs can  
use a combiner

Use a combiner

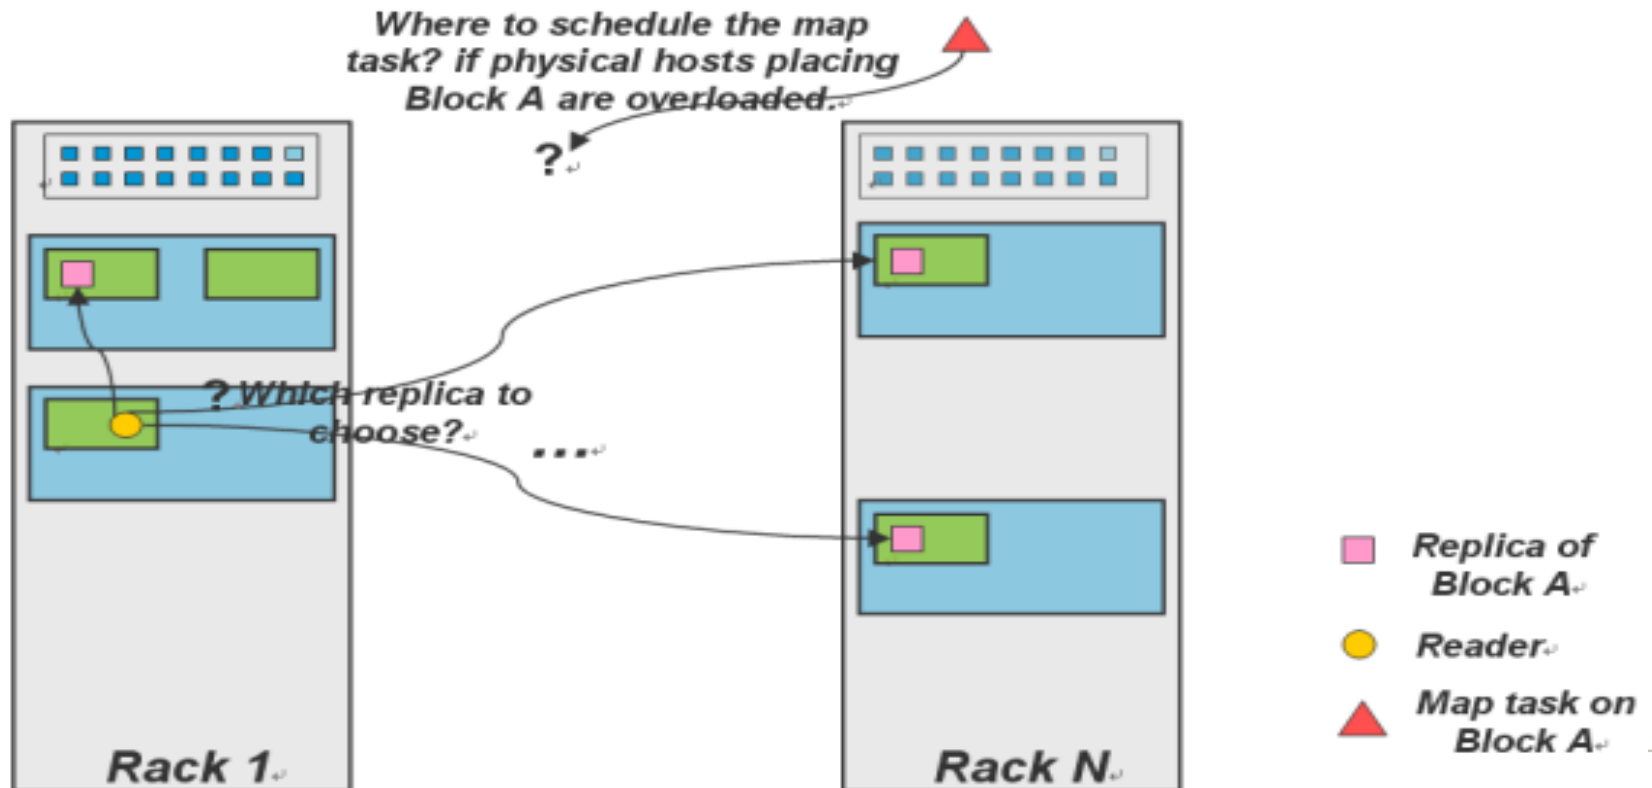
# Optimizations 3:

## Speculative Execution

- **Problem:** If one node is slow, it slows the entire job
- **Solution: Speculative Execution**
  - Hadoop automatically runs each task multiple times in parallel on different nodes
  - First one finishes, its result is used
  - Others will be killed

# Optimizations 4: Locality

- **Locality:** Run map code on same machine with relevant data
  - If not possible, then machine in the same rack
  - Best effort, as no guarantees could be given



# Optimizations

---

- **Optimizations for speed-up or scale-up of map-reduce processing**