

MongoDB:

Model and Operations

Data Model

Data Model

- BSON format (binary JSON)
- Developers can easily map to modern object-oriented languages without a complicated ORM layer.
- Lightweight, traversable, efficient

Terms : Relational DB vs. MongoDB

RDBMS	MongoDB
Database	Database
Table	Collection
Tuple/Row	Document
column	Field
Table Join	Embedded Documents
Primary Key	Primary Key (Default key _id provided by mongodb itself)
Database Server and Client	
Mysqld/Oracle	mongod
mysql/sqlplus	mongo

JSON

Field Name

Field Value

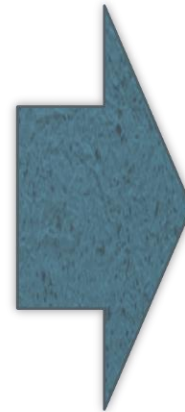
```
{
  "firstName": "John",
  "lastName": "Smith",
  "isAlive": true,
  "age": 25,
  "height_cm": 167.6,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null
}
```

- ***Field Value:***
 - Scalar (Int, Boolean, String, Date, ...)
 - Document (Embedding or Nesting)
 - Array of JSON objects

Example: JSON to BSON

```
{ author: 'joe',  
  created : new Date('03/28/2009'),  
  title : 'Yet another blog post',  
  text : 'Here is the text...',  
  tags : [ 'example', 'joe' ],  
  comments : [  
    { author: 'jim',  
      comment: 'I disagree'  
    },  
    { author: 'nancy',  
      comment: 'Good post'  
    }  
  ]  
}
```

For efficiency, stored in binary formats (BSON):



```
"\x16\x00\x00\x00\x02hello\x00  
\x06\x00\x00\x00world\x00\x00"  
  
"1\x00\x00\x00\x04BSON\x00&\x00  
\x00\x00\x020\x00\x08\x00\x00  
\x00awesome\x00\x011\x00333333  
\x14@\x102\x00\xc2\x07\x00\x00  
\x00\x00"
```

MongoDB Model

One **document** (e.g., one **tuple** in RDBMS)

```
{
  name: "sue",
  age: 26,
  status: "A",
  groups: [ "news", "sports" ]
}
```

← field: value
← field: value
← field: value
← field: value

- **Document** is equivalent to a tuple
- **Collection** groups similar documents

One **collection** (e.g., one **Table** in RDBMS)

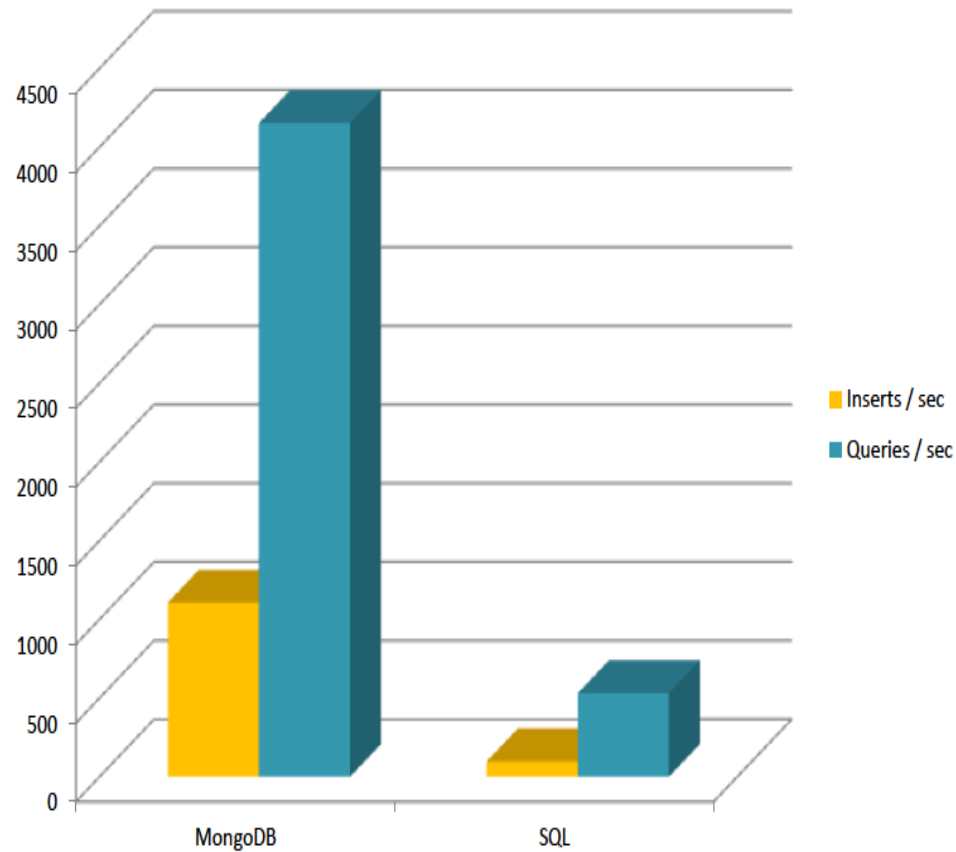
```
{
  name: "al",
  age: 18,
  status: "D",
  groups: [ "politics", "news" ]
}
```

Collection

- Within a collection, each document is unique

**Unlike RDBMS:
No Integrity
Constraints in
MongoDB**

Performance



**Unlike
RDBMS:
No Integrity
Constraints
in
MongoDB**

MongoDB Model (restrictions)

One *document* (e.g., one *tuple* in RDBMS)

```
{
  name: "sue",
  age: 26,
  status: "A",
  groups: [ "news", "sports" ]
}
```

← field: value
← field: value
← field: value
← field: value

One *Collection* (e.g., one *table* in RDBMS)

```
{
  name: "al",
  age: 18,
  status: "D",
  groups: [ "politics", "news" ]
}
```

Collection

Field names:

- **cannot** start with the **\$** character (command)
- **cannot** contain the **“.”** character (position)
- Max size of single document 16MB

```
{
  id: ObjectId(7df78ad8902c)
  title: 'MongoDB Overview',
  description: 'MongoDB is no sql database',
  by: 'tutorials point',
  url: 'http://www.tutorialspoint.com',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 100,
  comments: [
    {
      user: 'user1',
      message: 'My first comment',
      dateCreated: new Date(2011,1,20,2,15),
      like: 0
    },
    {
      user: 'user2',
      message: 'My second comments',
      dateCreated: new Date(2011,1,25,7,45),
      like: 5
    }
  ]
}
```

Identifier in MongoDB

- `_id` is special column in each document
- Unique within each collection
- `_id` \leftrightarrow Primary Key in RDBMS
- `_id` is 12 Bytes, you can set it yourself
- Or, system generated:
 - 1st 4 bytes \rightarrow timestamp
 - Next 3 bytes \rightarrow machine id
 - Next 2 bytes \rightarrow process id
 - Last 3 bytes \rightarrow incremental values

No Defined Schema (Schema-free Or Schema-less)

- MongoDB does not need any defined data schema.
- Every document could have different data!

```
{name: "will",  
  eyes: "blue",  
  birthplace: "NY",  
  aliases: ["bill", "la  
ciacco"],  
  gender: "???",  
  boss: "ben"}
```

```
{name: "jeff",  
  eyes: "blue",  
  height: 72,  
  boss: "ben"}
```

```
{name: "brendan",  
  aliases: ["el diablo"]}
```

```
{name: "ben",  
  hat: "yes"}
```

```
{name: "matt",  
  pizza: "DiGiorno",  
  height: 72,  
  boss: 555.555.1212}
```

Data Model Comparison:

Relational DB vs. NoSQL

This is hard...

Long time to develop

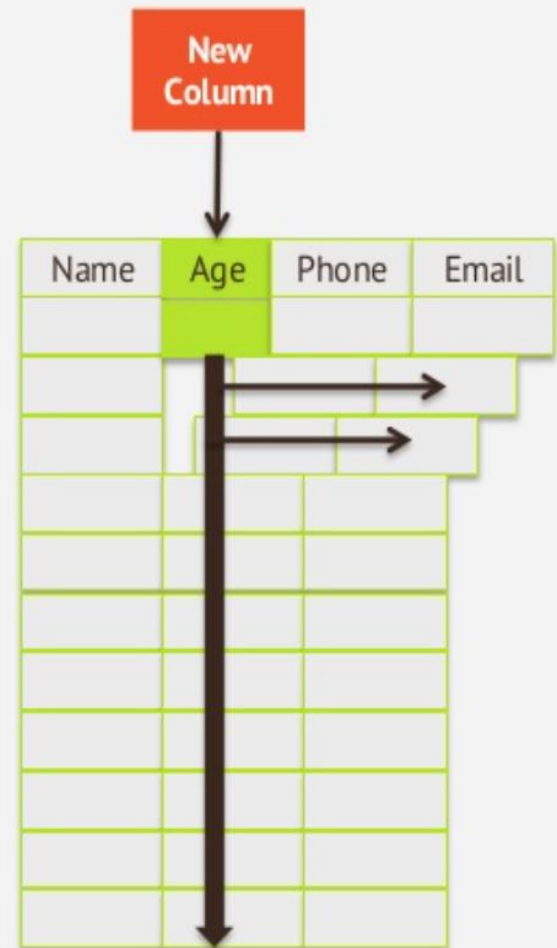
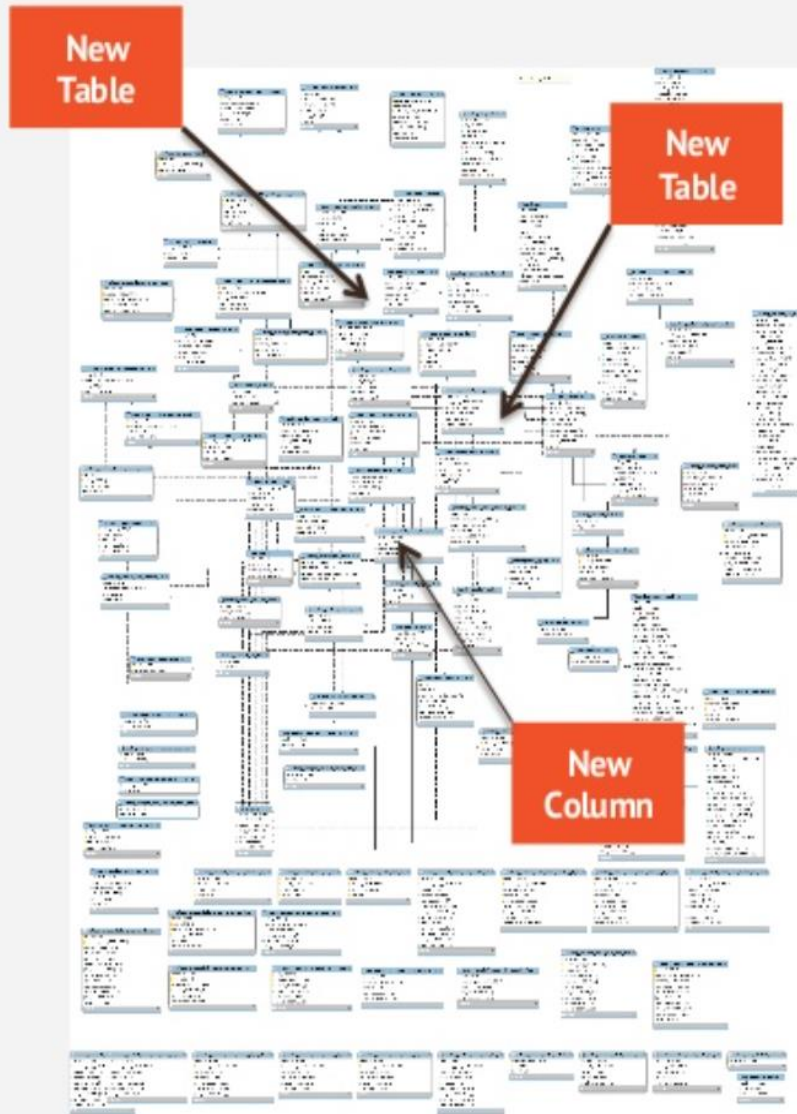
Difficult to change

- Complex relationships
- Dynamic environment

RDBMS are not always
the best choice

Queries are complex

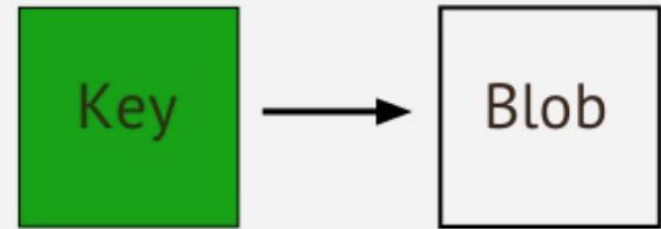
Hard to make changes



Key-Value Data Model (minimal)

Key → Value store

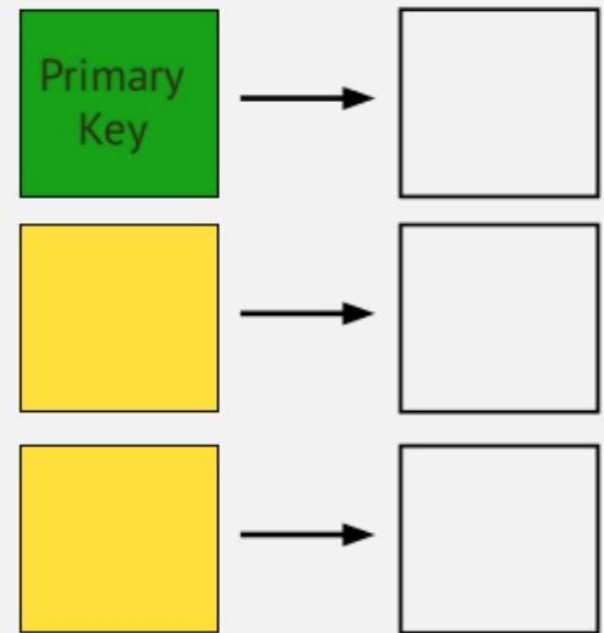
- One-dimensional storage
- Single value is a blob
- Query on key only
- No schema
- Value can be replaced but not updated



Relational Data Model (struct.)

Relational Record

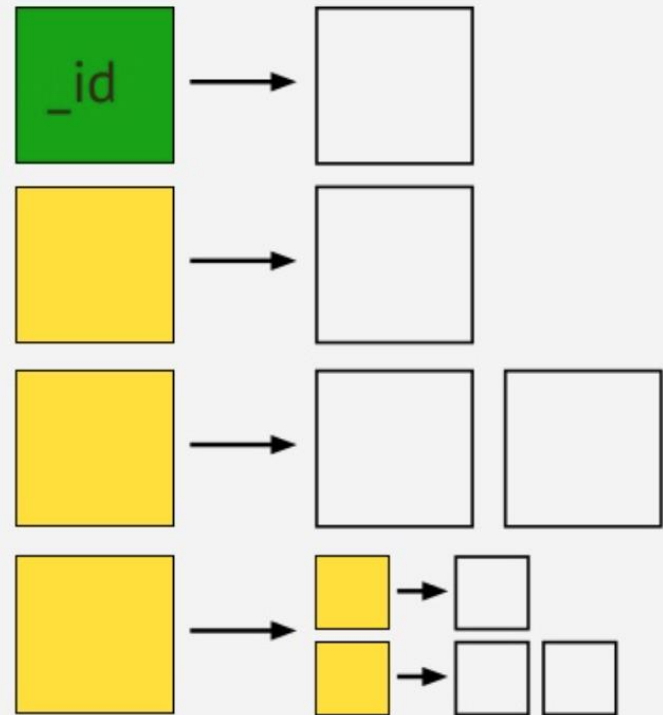
- Two-dimensional storage
- Field contains a single value
- Query on any field
- Very structured schema
- Poor data locality requires many tables, joins, and indexes.



Document Data Model (rich)

MongoDB Document

- **N-dimensional** storage
- Field can contain **many** values and **embedded** values
- Query on **any field & level**
- **Flexible** schema
- Optimal data locality requires fewer **indexes** and provides better **performance**

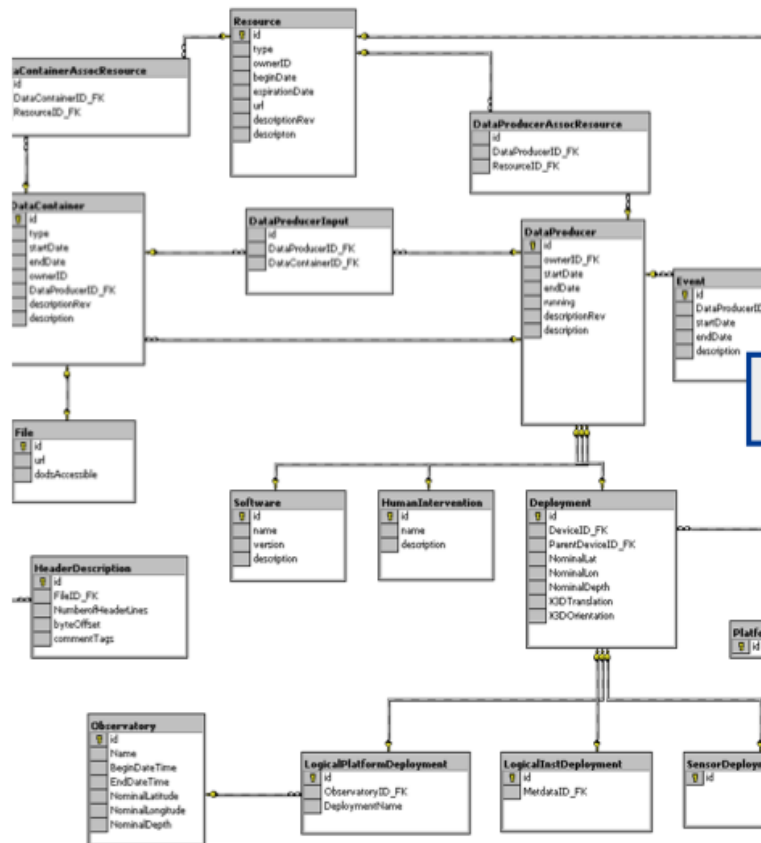


Document vs. Relational Models

- **Relational**
 - Focus on data storage
 - At query time → build your business objects
- **Document**
 - Focus on data usage
 - Maintain your business object in storage



Tradeoff: Normalization vs. Predetermined Usage



```
{  
  title: 'MongoDB',  
  contributors: [  
    { name: 'Eliot Horowitz',  
      email: 'eliot@10gen.com' },  
    { name: 'Dwight Merriman',  
      email: 'dwight@10gen.com' }  
  ],  
  model: {  
    relational: false,  
    awesome: true  
  }  
}
```

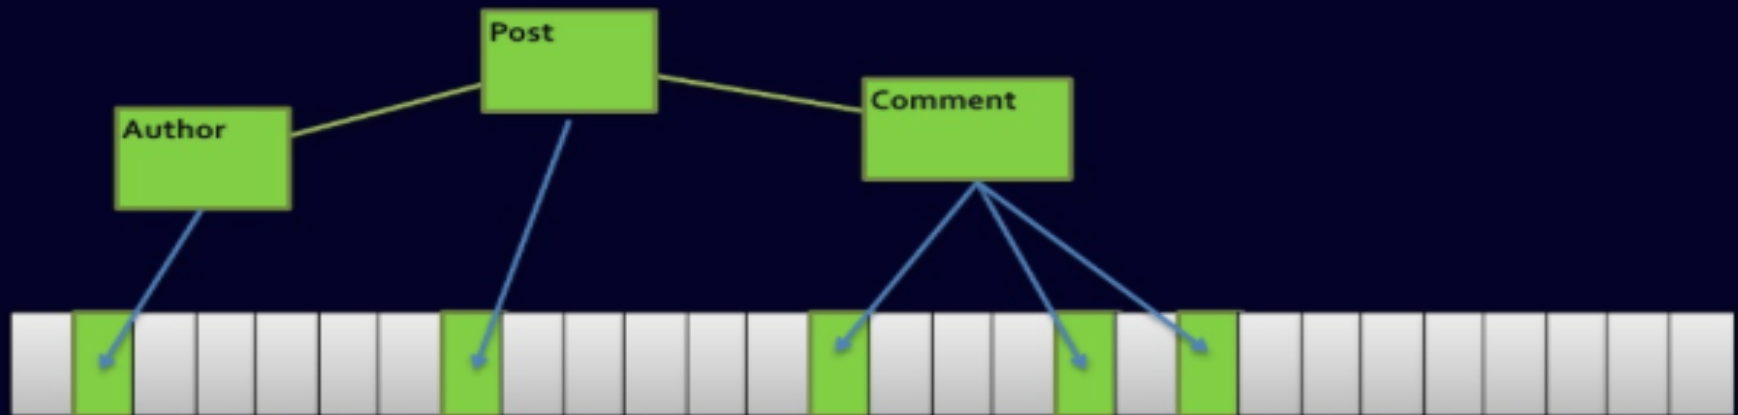
Complex Join Queries

Relational DBs

Disk seeks and data locality

Seek = 5+ ms

Read = really really fast



No Joins in MongoDB

Disk seeks and data locality



Updating & Querying

CRUD

- **C**reate
 - `db.collection.insert(<document>)`
 - `db.collection.update(<query>, <update>, { upsert: true })`
- **R**ead
 - `db.collection.find(<query>, <projection>)`
 - `db.collection.findOne(<query>, <projection>)`
- **U**ppdate
 - `db.collection.update(<query>, <update>, <options>)`
- **D**eleate
 - `db.collection.remove(<query>, <justOne>)`

CRUD Examples

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.find ()  
{  
  "_id" : ObjectId("51..."),  
  "first" : "John",  
  "last" : "Doe",  
  "age" : 39  
}
```

```
> db.user.update(  
  {"_id" : ObjectId("51...")},  
  {  
    $set: {  
      age: 40,  
      salary: 7000}  
  }  
)
```

```
> db.user.remove({  
  "first": /^J/  
})
```


Examples

In RDBMS

```
CREATE TABLE users (  
  id MEDIUMINT NOT NULL  
    AUTO_INCREMENT,  
  user_id Varchar(30),  
  age Number,  
  status char(1),  
  PRIMARY KEY (id)  
)
```

```
DROP TABLE users
```

In MongoDB

Create “Users” collection explicitly

```
db.createCollection("users")
```

```
db.users.drop()
```

Examples

In RDBMS

```
CREATE TABLE users (  
  id MEDIUMINT NOT NULL  
    AUTO_INCREMENT,  
  user_id Varchar(30),  
  age Number,  
  status char(1),  
  PRIMARY KEY (id)  
)
```

```
DROP TABLE users
```

In MongoDB

Either insert the 1st document

```
db.users.insert( {  
  user_id: "abc123",  
  age: 55,  
  status: "A"  
} )
```



Or create “Users” collection explicitly

```
db.createCollection("users")
```

```
db.users.drop()
```

Insertion (Basic Operation)

Collection
↓
`db.users.insert(`

Document
↓

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

)

Document

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

insert →

Collection

{ name: "al", age: 18, ... }
{ name: "lee", age: 28, ... }
{ name: "jan", age: 21, ... }
{ name: "kai", age: 38, ... }
{ name: "sam", age: 18, ... }
{ name: "mel", age: 38, ... }
{ name: "ryan", age: 31, ... }
{ name: "sue", age: 26, ... }

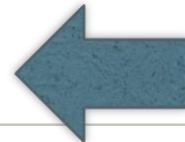
users

Collection "users" is created automatically if it does not exist

Multi-Document Insertion (Use of Arrays)

```
var mydocuments =  
  [  
    {  
      item: "ABC2",  
      details: { model: "14Q3", manufacturer: "M1 Corporation" },  
      stock: [ { size: "M", qty: 50 } ],  
      category: "clothing"  
    },  
    {  
      item: "MNO2",  
      details: { model: "14Q3", manufacturer: "ABC Company" },  
      stock: [ { size: "S", qty: 5 }, { size: "M", qty: 5 }, { size: "L", qty: 1 } ],  
      category: "clothing"  
    },  
    {  
      item: "IJK2",  
      details: { model: "14Q2", manufacturer: "M5 Corporation" },  
      stock: [ { size: "S", qty: 5 }, { size: "L", qty: 1 } ],  
      category: "houseware"  
    }  
  ];
```

```
db.inventory.insert( mydocuments );
```



All documents are
inserted at once into
inventory

Multi-Document Insertion (Bulk Operation)

- Temporary object in memory to hold your insertions and upload them at once into storage

There is a *Bulk Ordered* object



```
var bulk = db.inventory.initializeUnorderedBulkOp();
bulk.insert(
  {
    item: "BE10",
    details: { model: "14Q2", manufacturer: "XYZ Company" },
    stock: [ { size: "L", qty: 5 } ],
    category: "clothing"
  }
);
bulk.insert(
  {
    item: "ZYT1",
    details: { model: "14Q1", manufacturer: "ABC Company" },
    stock: [ { size: "S", qty: 5 }, { size: "M", qty: 5 } ],
    category: "houseware"
  }
);

bulk.execute();
```

_id column is added
automatically

Deletion (Remove Operation)

- You can put condition on any field in the document (even *id*)

```
db.users.remove(  
  { status: "D" }  
)
```

← collection
← remove criteria

The following diagram shows the same query in SQL:

```
DELETE FROM users  
WHERE status = 'D'
```

← table
← delete criteria

db.users.remove ()




Removes all documents from *users* collection

Update

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option



Otherwise, it will update only 1st matching document

Equivalent to SQL:

```
UPDATE users  
SET      status = 'A'  
WHERE    age > 18
```

← table
← update action
← update criteria

Update (Cont'd)

Two update operators

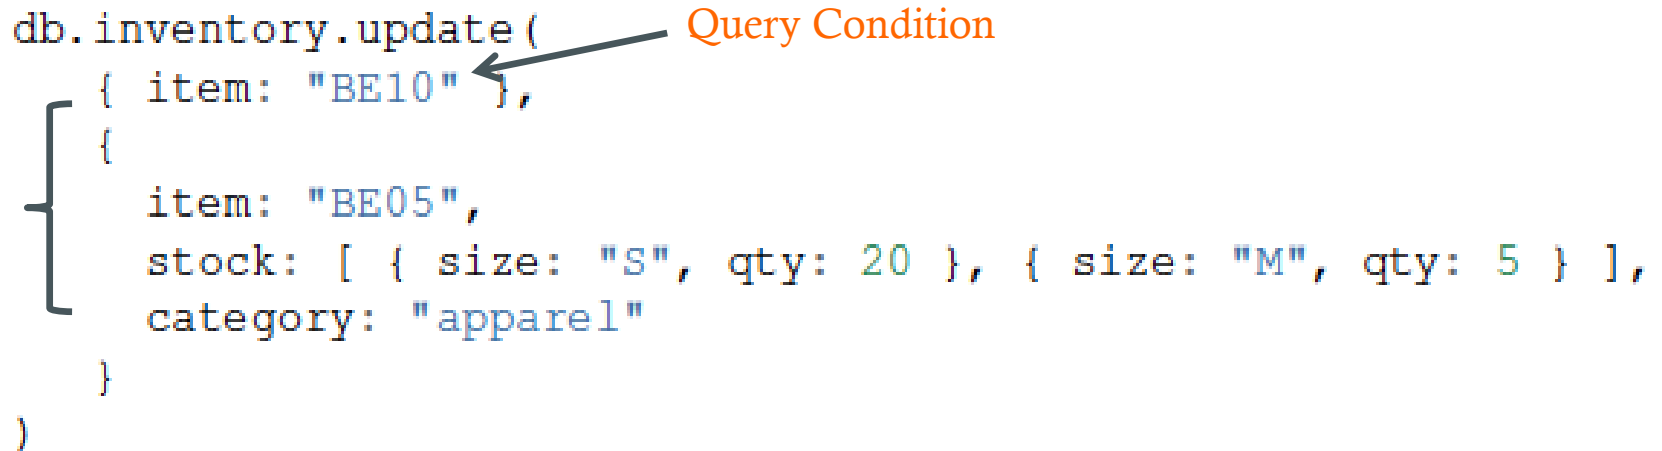
```
db.inventory.update(  
  { item: "MNO2" },  
  {  
    $set: {  
      category: "apparel",  
      details: { model: "14Q3", manufacturer: "XYZ Company" }  
    },  
    $currentDate: { lastModified: true }  
  }  
)
```

For the document with `item` equal to "MNO2", use the `$set` operator to update the `category` field and the `details` field to the specified values and the `$currentDate` operator to update the field `lastModified` with the current date.

Replace a document

New
doc

```
db.inventory.update(  
  { item: "BE10" },  
  {  
    item: "BE05",  
    stock: [ { size: "S", qty: 20 }, { size: "M", qty: 5 } ],  
    category: "apparel"  
  }  
)
```

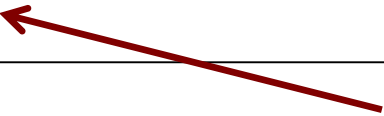


Because there is no \$SET operator;

We replace the document having item = “BE10”, with the given document

Insert or Replace

```
db.inventory.update(  
  { item: "TBD1" },  
  {  
    item: "TBD1",  
    details: { "model" : "14Q4", "manufacturer" : "ABC Company" },  
    stock: [ { "size" : "S", "qty" : 25 } ],  
    category: "houseware"  
  },  
  { upsert: true }  
)
```



The *upsert* option

**If the document having item = “TBD1” is in the DB, it will be replaced
Otherwise, it will be inserted.**

REMINDER

Installation: in your VM

Manual: <https://docs.mongodb.com/manual>

Data examples:

<https://docs.mongodb.com/manual/reference/bios-example-collection/>

Online Terminal:

http://www.tutorialspoint.com/mongodb_terminal_online.php