

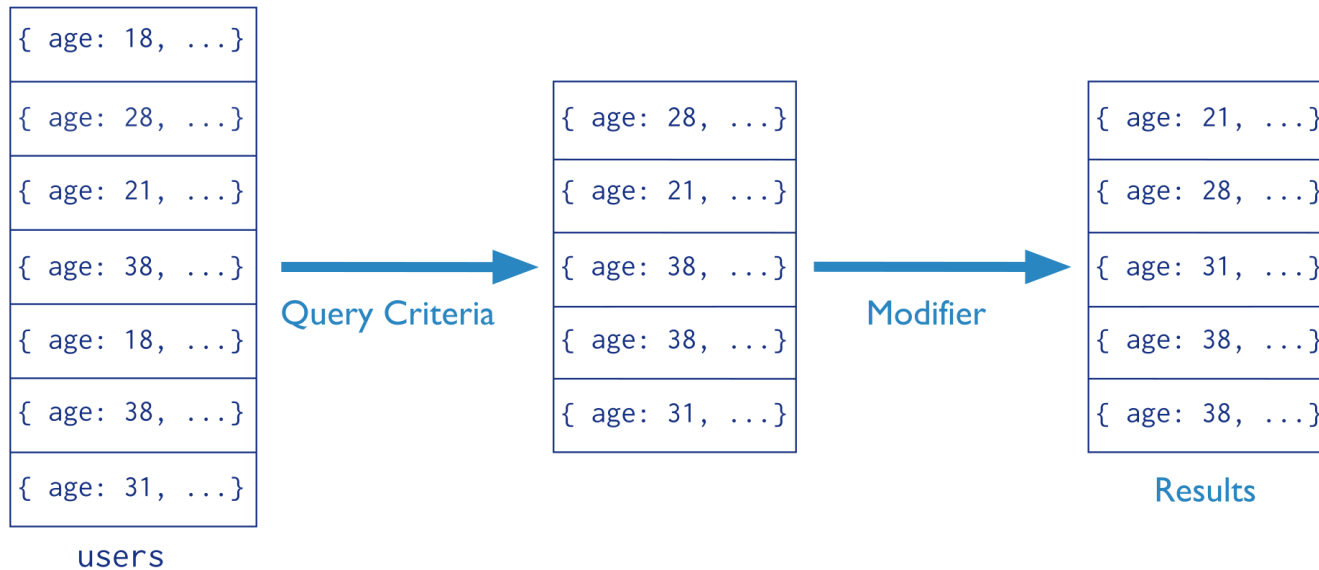
# Query Language in MongoDB

# Find() Operator

Collection      Query Criteria      Modifier

db.users.find( { age: { \$gt: 18 } } ).sort( {age: 1 } )

Means ascending



# Find() + Projection

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
)
```



← collection  
← query criteria  
← projection  
← cursor modifier

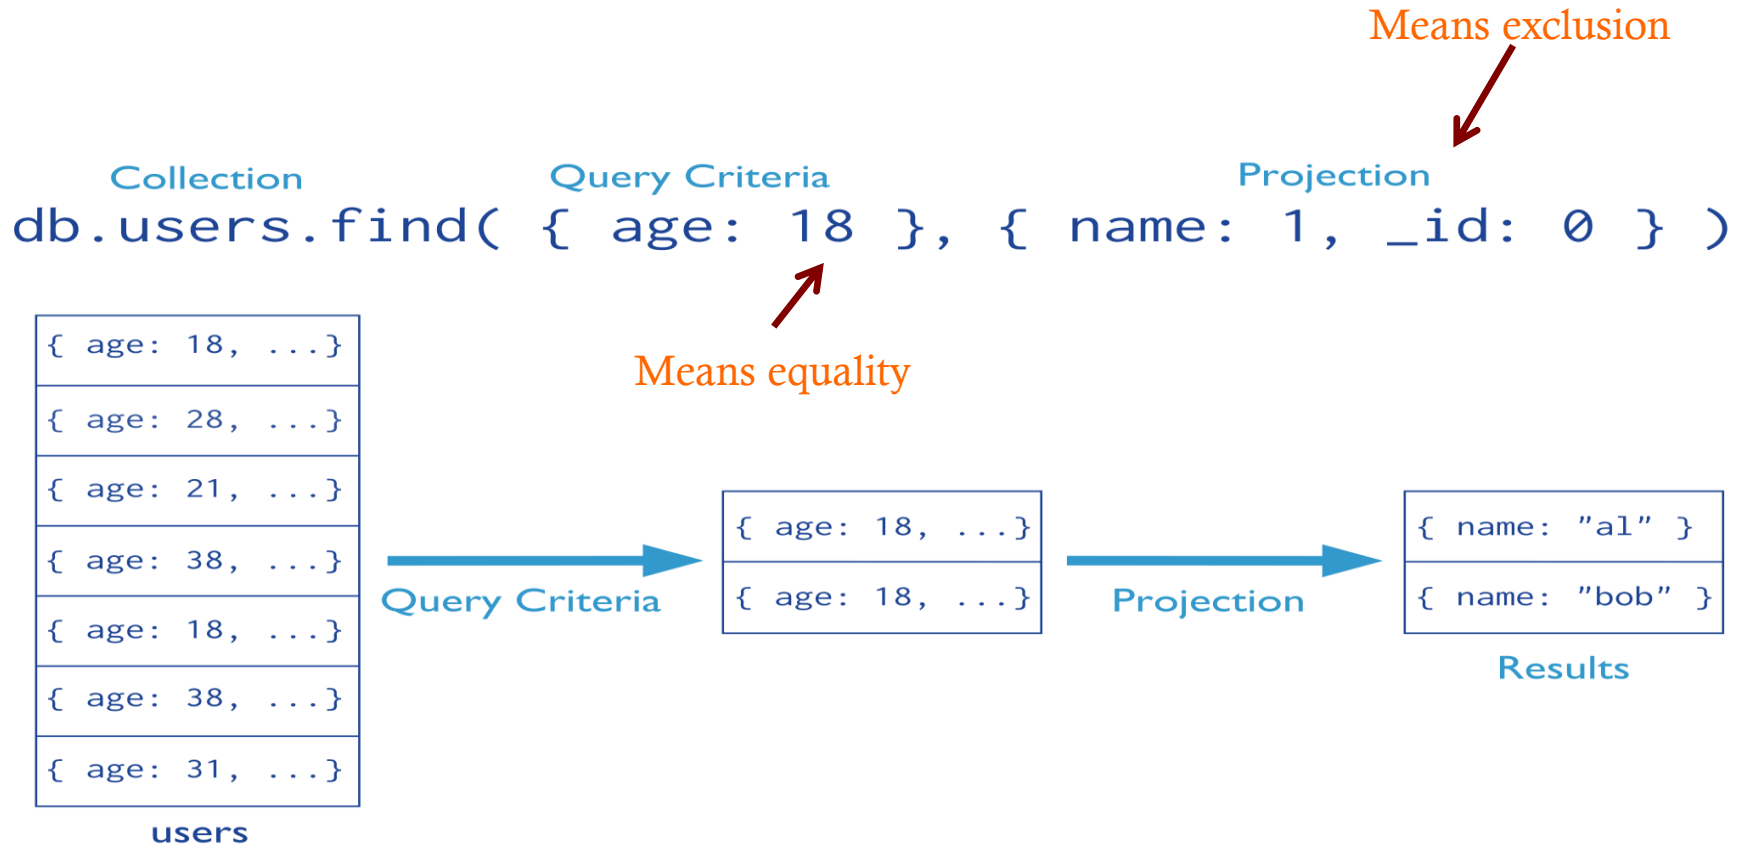
Means inclusion +  
***\_id is always automatically included***

**Equivalent to in SQL:**

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

← projection  
← table  
← select criteria  
← cursor modifier

# Find(): Exclude Fields



Cannot mix “inclusion & exclusion” in the same operator except for *id*

# More Examples for Find()

Report all documents in the “inventory” collection

```
db.inventory.find( )
```

**Equivalent to in SQL:**

```
Select *  
From inventory;
```

Report all documents in the “inventory” collection  
Where **type = ‘food’ or ‘snacks’**

```
db.inventory.find(  
  { type: { $in: [ 'food', 'snacks' ] } }  
)
```

**Equivalent to in SQL:**

```
Select *  
From inventory  
Where type in  
  ('food', 'snacks');
```

# Find(): AND & OR Connectors

## AND Semantics (default and thus implicit connector)

```
db.inventory.find( { type: 'food', price: { $lt: 9.95 } } )
```

## OR Semantics

```
db.inventory.find(  
  {  
    $or: [ { qty: { $gt: 100 } }, { price: { $lt: 9.95 } } ]  
  }  
)
```

## AND + OR Semantics

```
db.inventory.find(  
  {  
    type: 'food',  
    $or: [ { qty: { $gt: 100 } }, { price: { $lt: 9.95 } } ]  
  }  
)
```

Type = 'food' and (qty > 100 or price < 9.95)

# \$AND: Logical AND Operation

```
{ $and: [ { <expression1> }, { <expression2> } , ... , {<expressionN> } ] }
```

```
db.inventory.find(  
  { $and :  
    [  
      { $or : [ { price : 0.99 }, { price : 1.99 } ] },  
      { $or : [ { sale : true }, { qty : { $lt : 20 } } ] }  
    ]  
  } )
```

# Queries Return Cursors

- All queries return the results in a cursor
- If not assigned to a variable → Printed to screen
  - Results are stored in a cursor
  - Operators on results can manipulate the cursor

**var** myCursor = db.users.find( { type: 2 } );

**while** (myCursor.hasNext())

{ print( tojson(myCursor.next())); }

**Cursor's Methods:** <http://docs.mongodb.org/manual/reference/method/js-cursor/>




# Querying Complex Types

# Querying Complex Types

```
{
  "firstName": "John",
  "lastName": "Smith",
  "isAlive": true,
  "age": 25,
  "height_cm": 167.6,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null
}
```

Documents can be complex, e.g., arrays, embedded documents, any nesting of these, many levels



Queries can query these complex types

# Array Manipulation (Exact Match)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: [ 5, 8, 9 ] } )
```

# Array Manipulation (Exact Match)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: [ 5, 8, 9 ] } )
```

The operation returns the following document:

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }
```

# Array Manipulation (Search By Element)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: 5 } )
```



# Array Manipulation (Search By Element)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: 5 } )
```

The operation returns the following documents:

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 6, "type" : "food", "item" : "bbb", "ratings" : [ 5, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

**Notice:** if a document has “ratings” as an Integer field = 5, it will be returned

# Array Manipulation (Search By Position)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { 'ratings.0': 5 } )
```



# Array Manipulation (Search By Position)

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { 'ratings.0': 5 } )
```

The operation returns the following documents:

```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 6, "type" : "food", "item" : "bbb", "ratings" : [ 5, 9 ] }
```



# Array Manipulation (\$elemMatch)


```
{ <field>: { $elemMatch: { <query1>, <query2>, ... } } }
```

\$elemMatch operator matches documents that contain an **array** field with at least **one** element that matches **all** specified query criteria.

# Another Example

```
{ _id: 5, type: "food", item: "aaa", ratings: [ 5, 8, 9 ] }  
{ _id: 6, type: "food", item: "bbb", ratings: [ 5, 9 ] }  
{ _id: 7, type: "food", item: "ccc", ratings: [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: { $elemMatch: { $gt: 5, $lt: 9 } } } )
```



```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

```
db.inventory.find( { ratings: { $gt: 5, $lt: 9 } } )
```



```
{ "_id" : 5, "type" : "food", "item" : "aaa", "ratings" : [ 5, 8, 9 ] }  
{ "_id" : 6, "type" : "food", "item" : "bbb", "ratings" : [ 5, 9 ] }  
{ "_id" : 7, "type" : "food", "item" : "ccc", "ratings" : [ 9, 5, 8 ] }
```

# Embedded Object Matching (Exact document Matching)

```
{  
  name: "Joe",  
  address: {  
    city: "San Francisco",  
    state: "CA" },  
  likes: [ 'scuba', 'math', 'literature' ]  
}
```

```
db.persons.find( { "address" : { state: "CA" } } ) //don't match
```

```
db.persons.find( { "address" : {city: "San Francisco", state: "CA" } } ) // match
```

```
db.persons.find( { "address" : {state: "CA" , city: "San Francisco"}} ) //don't match
```

Exact-match  
(entire object)

# Embedded Object Matching (Field Matching)

```
{  
  name: "Joe",  
  address: {  
    city: "San Francisco",  
    state: "CA" },  
  likes: [ 'scuba', 'math', 'literature' ]  
}
```

Find the user documents where the  
address's state = 'CA'

```
db.persons.find( {"address.state" : "CA"})
```



Using dot notation

# Matching Arrays of Embedded Documents

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```


Q: Select all documents where the memos array contains in the 1<sup>st</sup> element a document written by 'shipping' department

# Matching Arrays of Embedded Documents

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```

`db.inventory.find( { 'memos.0.by': 'shipping' })` //Returns 1<sup>st</sup> document



Means the 1<sup>st</sup> element in the array

# Matching Arrays of Embedded Documents

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```

Q: Select all documents where the memos array contains a document written “by” shipping department

# Matching Arrays of Embedded Documents

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```

`db.inventory.find( { 'memos.by': 'shipping' } )`     *// Returns both documents*

 Means any element in the array



# Matching Arrays of Embedded Documents: Multiple Conditions

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```

Q: Select all documents where memos array contains a **document** written by 'shipping' department and its content is “on time”

# Matching Arrays of Embedded Documents: Multiple Conditions

```
db.inventory.find(  
  {  
    memos:  
      {  
        $elemMatch:  
          {  
            memo: 'on time',  
            by: 'shipping'  
          }  
        }  
      }  
    }  
  )
```

# Matching Arrays of Embedded Documents: Multiple Conditions

```
{
  _id: 100,
  type: "food",
  item: "xyz",
  qty: 25,
  price: 2.5,
  ratings: [ 5, 8, 9 ],
  memos: [ { memo: "on time", by: "shipping" }, { memo: "approved", by: "billing" } ]
}

{
  _id: 101,
  type: "fruit",
  item: "jkl",
  qty: 10,
  price: 4.25,
  ratings: [ 5, 9 ],
  memos: [ { memo: "on time", by: "payment" }, { memo: "delayed", by: "shipping" } ]
}
```

```
db.inventory.find(
  {
    memos:
      {
        $elemMatch:
          {
            memo: 'on time',
            by: 'shipping'
          }
      }
  }
)
```

# Summary: Query Operators

- <http://docs.mongodb.org/manual/reference/operator/query/>

- Comparison Operators
- Logical Operators
- Element Operators
- Evaluation Operators
- Array Operators

# Query Operators: Comparison Op

Name	Description
<code>\$eq</code>	Matches values that are equal to a specified value.
<code>\$gt</code>	Matches values that are greater than a specified value.
<code>\$gte</code>	Matches values that are greater than or equal to a specified value.
<code>\$lt</code>	Matches values that are less than a specified value.
<code>\$lte</code>	Matches values that are less than or equal to a specified value.
<code>\$ne</code>	Matches all values that are not equal to a specified value.
<code>\$in</code>	Matches any of the values specified in an array.
<code>\$nin</code>	Matches none of the values specified in an array.

```
db.inventory.find  
( { qty: { $gte: 20 } } )
```

```
db.inventory.update(  
  { "carrier.fee": { $gte: 2 } },  
  { $set: { price: 9.99 } }  
)
```

# Query Operators: Evaluation Op

Name	Description
<code>\$mod</code>	Performs a modulo operation on the value of a field and selects documents with a specified result.
<code>\$regex</code>	Selects documents where values match a specified regular expression.
<code>\$text</code>	Performs text search.
<code>\$where</code>	Matches documents that satisfy a JavaScript expression.

# \$Where Operator

- Passes a *JavaScript expression or function* to query system
- Very flexible in expressing complex conditions
- But relatively slow as it evaluates for each document (no index)
- Similar to using *UDF* in *WHERE* clause in relational database

```
db.myCollection.find( { $where: "this.credits == this.debits" } );
```

```
db.myCollection.find( { $where: "obj.credits == obj.debits" } );
```

```
db.myCollection.find( { $where: function() { return (this.credits == this.debits) } } );
```

```
db.myCollection.find( { $where: function() { return obj.credits == obj.debits; } } );
```

# \$Where Operator

- Can combine MongoDB operators with \$Where

```
db.myCollection.find(  
  { active: true, $where: "this.credits - this.debits < 0" } );
```

```
db.myCollection.find(  
  { active: true,  
    $where: function() { return obj.credits - obj.debits < 0; } }  
);
```



# Collection Modeling

# Collection Modeling

- Modeling multiple collections that reference each other
- In Relational DBs ➔ FK-PK Relationships
- In MongoDB, two options:
  - Referencing
  - Embedding

# FK-PK in Relational DBs

- Create “Students” relation

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa REAL);
```

- Create “Courses” relation

```
CREATE TABLE Courses
(cid Varchar2(20),
 name varchar2(50),
 maxCredits integer,
 graduateFlag char(1));
```

Foreign key

Foreign key

- Create “Enrolled” relation

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid Varchar2(20),
 enrollDate date,
 grade CHAR(2));
```

◆ Each tuple in  
“Enrolled” reference a  
specific student and a  
specific course

# FK-PK in Relational DBs

## It comes with an enforcement mechanism

- Cannot insert a FK for a non-existing PK
- You cannot delete a PK that has a FK

**Enrolled** (referencing relation)

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Foreign Key

**Students** (referenced relation)

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Primary Key

# “Relationships” In MongoDB

- *Referencing* between two collections
  - Use Id of one and put in the other collection
  - Very similar to FK-PK in Relational Databases
  - **HOWEVER: Does not come with any enforcement !**
- *Embedding* between two collections
  - Put the document from one collection inside the other one



# Modeling using Referencing

*No Enforcements*

user document

```
{
  _id: <ObjectId1>,
  username: "123xyz"
}
```

contact document

```
{
  _id: <ObjectId2>,
  user_id: <ObjectId1>,
  phone: "123-456-7890",
  email: "xyz@example.com"
}
```

access document

```
{
  _id: <ObjectId3>,
  user_id: <ObjectId1>,
  level: 5,
  group: "dev"
}
```

*“Normalized” Model*

- Have three collections in the DB: “User”, “Contact”, “Access”
- Link them by `_id` (or any other field(s))

# Embedding

*De-Normalized Way*

```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

- Have one collection in DB as the “User” document
- The others are embedded inside a user’s document model

# Examples (1)

*Referencing*

- “Patron” & “Addresses”

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}
```

```
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketon",  
  state: "MA",  
  zip: "12345"  
}
```

- If it is 1-1 relationship
- If usually read the address with the name
- If address document usually does not expand

**If most of these hold  
→  
better use Embedding**



# Examples (2)

*Embedding*

- “Patron” & “Addresses”

```
{  
  _id: "joe",  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake Street",  
    city: "Faketon",  
    state: "MA",  
    zip: "12345"  
  }  
}
```

Advantages:

1. When you read, you get entire document at once
2. In Referencing → Need to issue multiple queries

# Examples (3)

*Referencing*

- What if a “Patron” can have many “Addresses” ?

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}
```

```
{  
  patron_id: "joe",  
  addresses:  
  {  
    patron_id: "joe",  
    street: "123 Fake Street",  
    city:  
    state:  
    zip:  
  }  
}
```

```
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketon",  
  state: "MA",  
  zip: "12345"  
}
```

- Do you read them together → Go for Embedding
- Are addresses dynamic (e.g., add new ones frequently)

→ Go for Referencing

# Examples (4) *Embedding*

If a “Patron” can have many “Addresses”, how to embed?

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

Use array of addresses

# Examples (5)

- If addresses are added frequently ... ?

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

This array will expand frequently



Size of "Patron" document increases frequently



May trigger re-locating the document each time

***Bad !!!!!***

# Document Size and Storage

- Each document contiguous on disk
- If doc size increases
  - ➔ Document location must change
- If doc location changes
  - ➔ Indexes must be updated
  - ➔ leads to expensive updates

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

## In a newer version of MongoDB:

Each document is allocated a *power-of-2 bytes* (smallest above its size)

System keeps some space empty for possible expansion

# More Examples (6)

- **One-to-Many “Book”, “Publisher”**

- A book has one publisher
- A publisher publishes many books

*How Model ?  
Pros and Cons ?*

- **If embed “Publisher” inside “Book”**

- Repeating publisher info inside each of its books
- Very hard to update publisher’s info (replicas, inconsistency!)

- **If embed “Book” inside “Publisher”**

- Book becomes an array (many)
- Frequent updates and increases in size (expensive)

# More Examples (6)

- **One-to-Many “Book”, “Publisher”**

- A book has one publisher
- A publisher publishes many books

*Referencing is better  
in this case*

- **If embed “Publisher” inside “Book”**

- Repeating publisher info inside each of its books
- Very hard to update publisher’s info

- **If embed “Book” inside “Publisher”**

- Book becomes an array (many)
- Frequently update and increases in size

# Summary:

Query Language in MongoDB:  
**Powerful document traversal  
and search operations**