



# Apache Spark

## Part 2

**MATERIALS ADOPTED FROM**  
**Berkeley's AMPLab ([Matei Zaharia](#) – PhD student in 2009),**  
**open-sourced in 2010 as BSD,**  
**APACHE software foundation in 2013 as SPARK**  
**Company formed DATABRICKS Inc! record in fast sorting**  
**in 2014.**

# *Generic Efficient* Infrastructure

Spark  
SQL

Spark  
Streaming

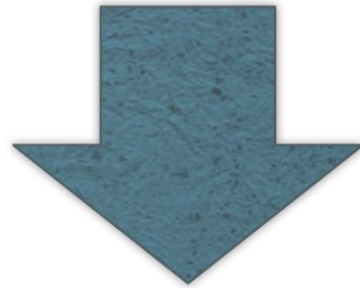
MLlib  
(machine  
learning)

GraphX  
(graph)

Apache Spark

# Motivation : Workloads

- Complex multi-pass algorithms
- Interactive ad-hoc queries
- Real-time stream processing

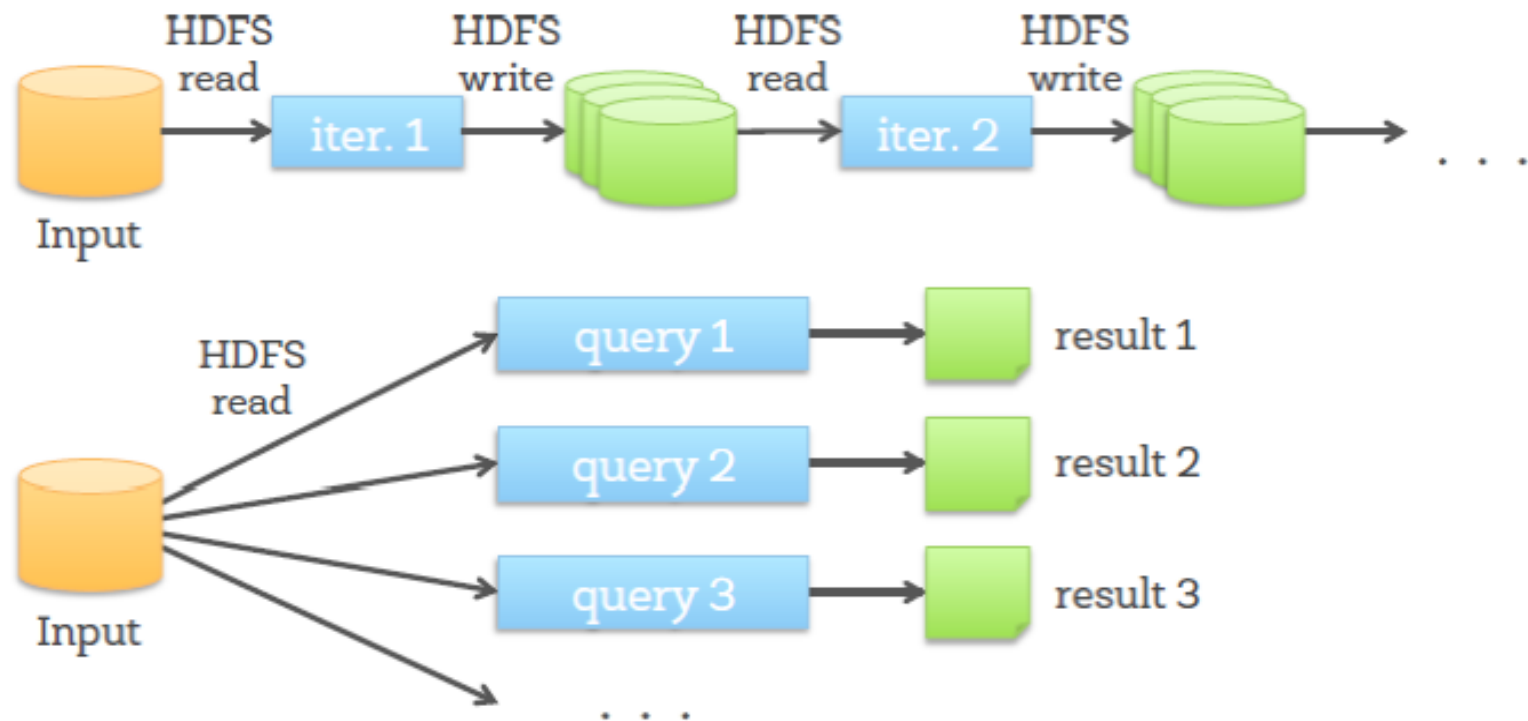


**All need efficient data sharing and transfer**

# Motivation: Serving Workloads

## From This ...

### Data Sharing in MapReduce

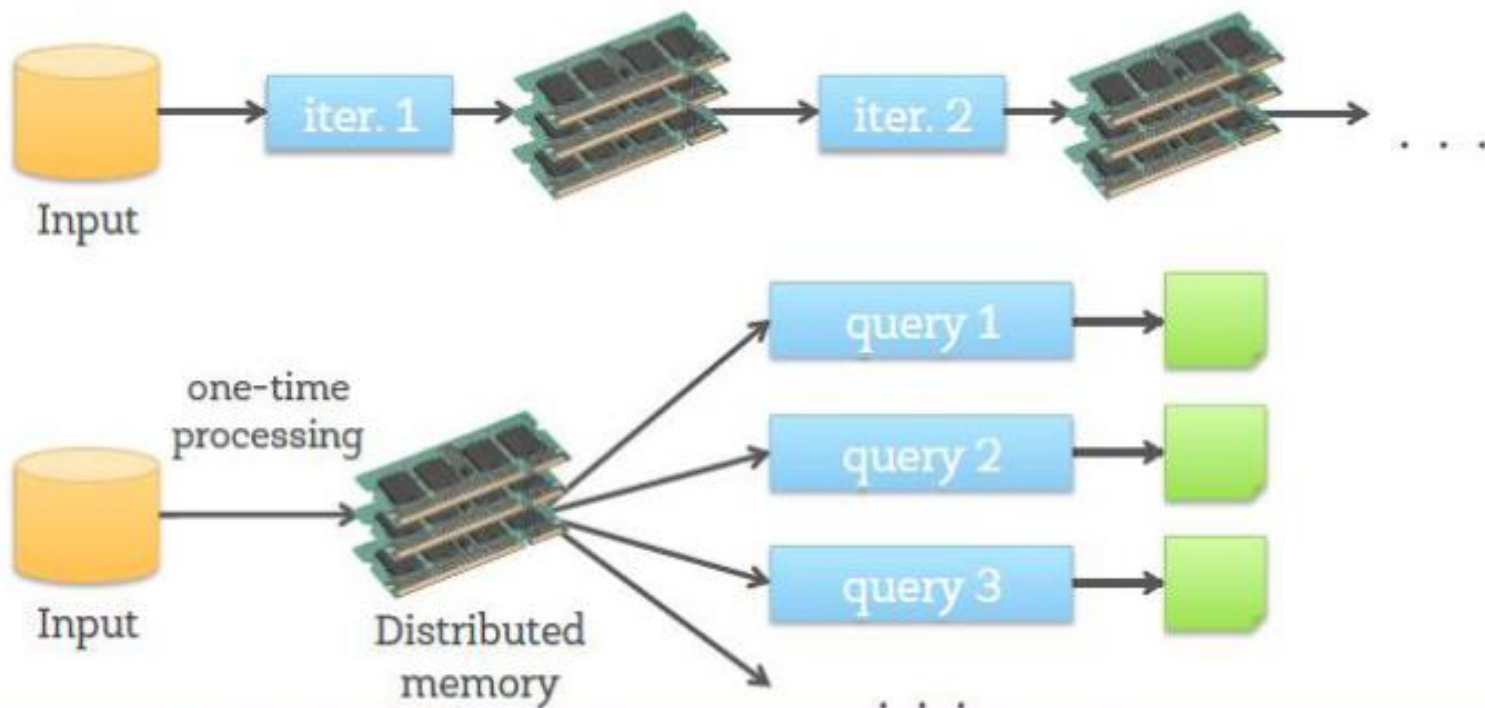


Slow due to data replication and disk I/O

# Motivation Workloads

## Thus To This ...

What We'd Like



10-100× faster than network and disk

# Motivation

## From Hardware Side

- RAM is getting much cheaper
- Commodity machines with GBs of RAM
- Large Distributed RAM in the cluster



**Processing, storage,  
and data transfer to  
use RAM,  
If possible.**

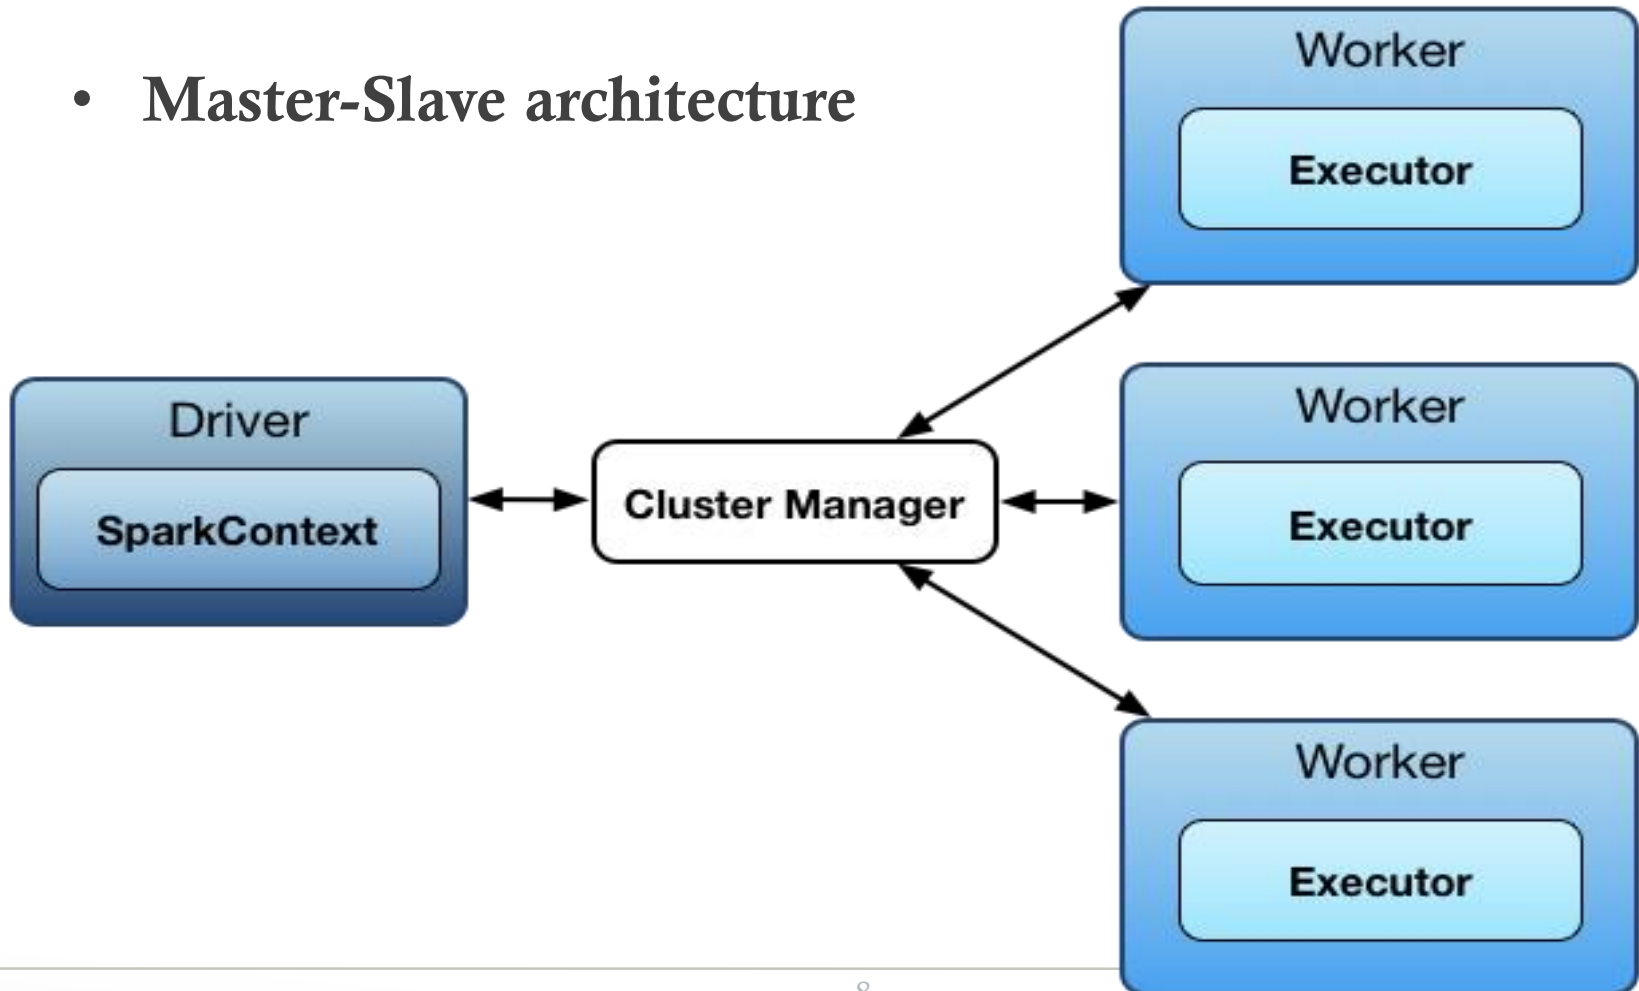
# Motivation: Overall

---

- Better support for real-time processing
- Exploit RAM as much as possible
- Large-scale distributed computations

# Spark Architecture

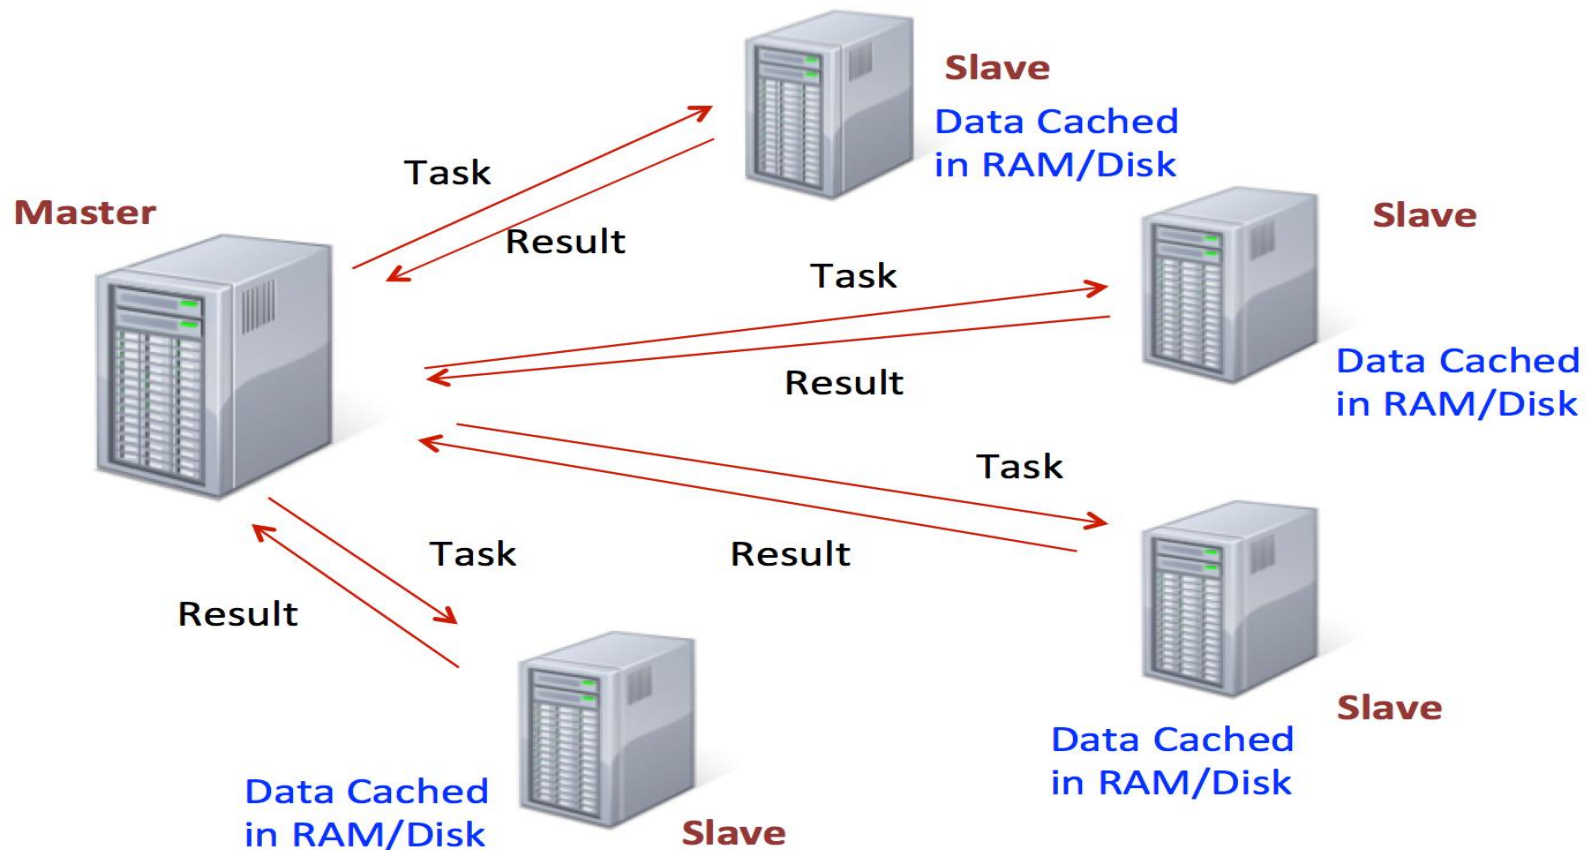
- Master-Slave architecture





# Spark Communication Model

## How does Spark execute a job



# Spark Programming Model

- High-level coding to build a workflow:  
**SCALA** as functional prog. language
- Code compiles to distributed parallel operations
- Two Abstraction Units:
  - **RDDs: Resilient Distributed Datasets**
  - **Paradigm: Parallel Operations**



- General purpose programming language (type-safe)
- Combines Object-Oriented and Functional programming
- Features: Concise, logical, and powerful language.
- Compiles to Java bytecode
- Runs on JVM

# Comparison

Java	Scala
Complex syntax	Simple syntax
Requires lengthy code.	Shorter code for same purpose.
Rewriting is needed.	Rewriting is not required.
Dynamic in nature	Statically-typed
No assurance of bug-free code	Assurance of lesser defects

# Spark RDDs

# RDD: Concept

## **Resilient Distributed Datasets :**

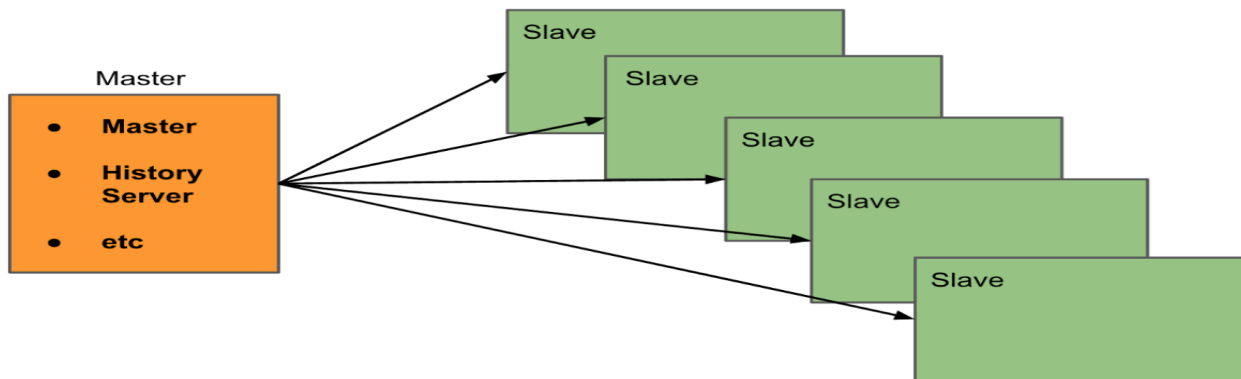
- “Collection of objects” (records) that act as one unit
- Stored in main memory (or when needed on disk)
- Parallel operations on top of this “collection”
- Have fault tolerance without replication (lineage)

# RDD: Concept

- RDD is read-only

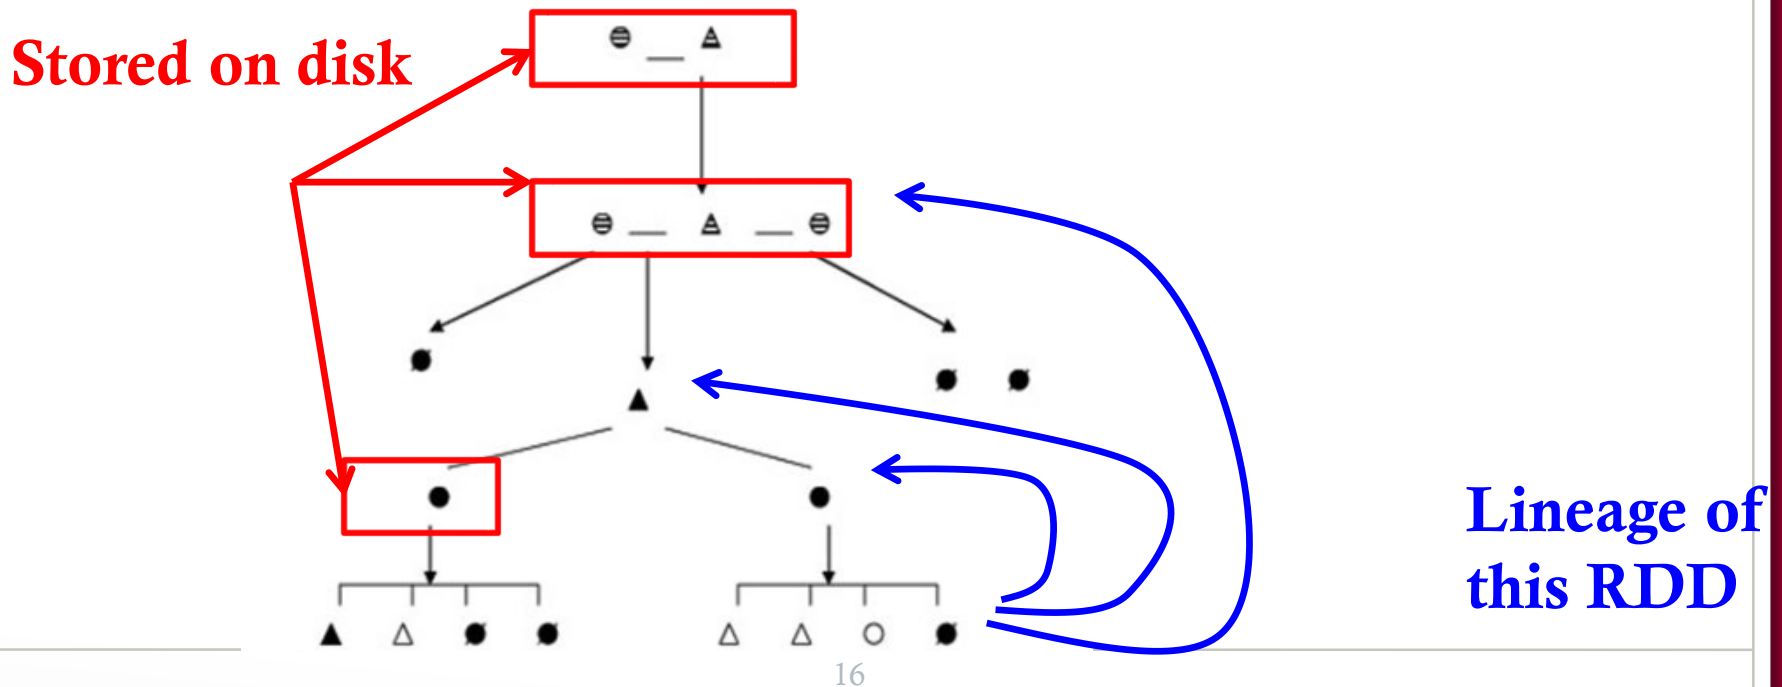


- Distributed either in main memory or disk (automatically decided)



# RDD: Fault Tolerance

- Not (fully) replicated. But maintain **lineage** (provenance) on how to re-create RDD starting from data in **reliable storage**

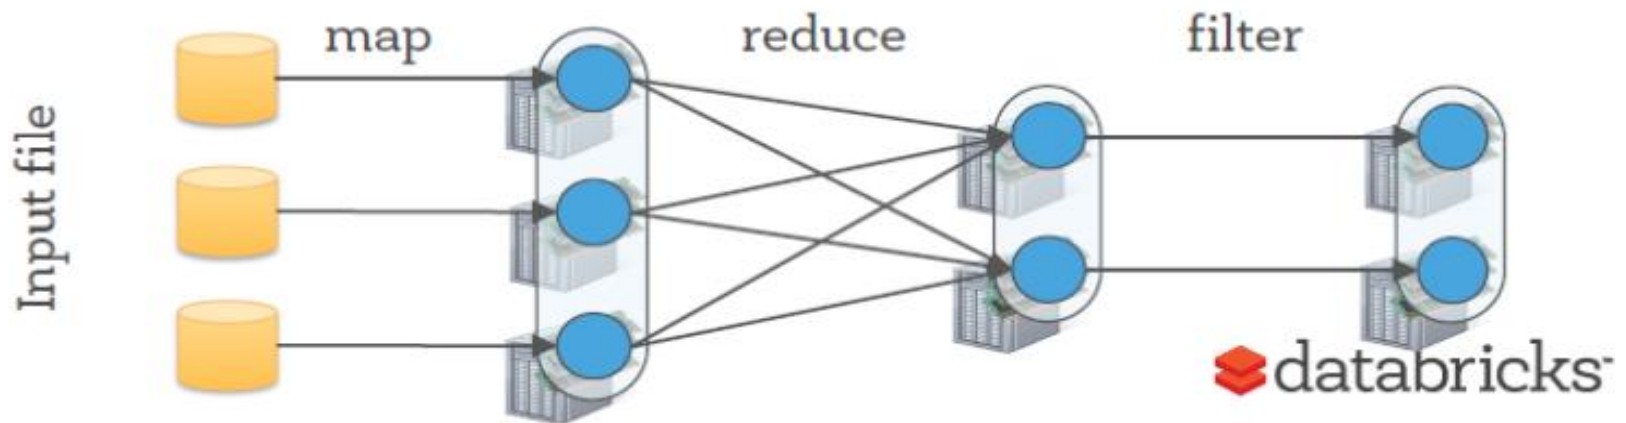




# RDD: Fault Tolerance

RDDs track *lineage* info to rebuild lost data

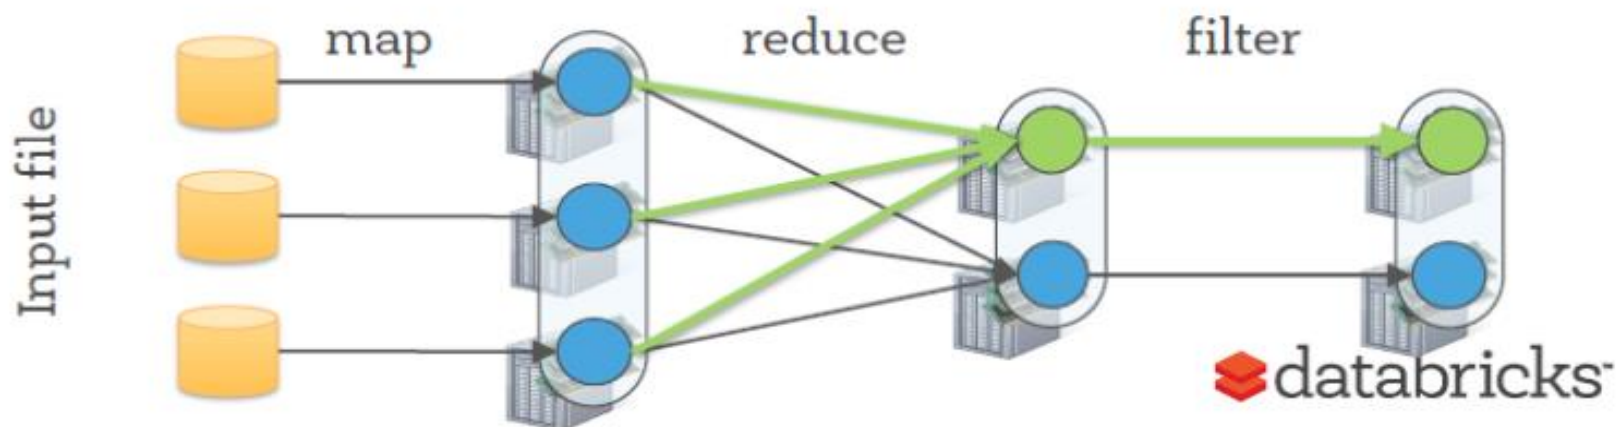
```
file.map(lambda rec: (rec.type, 1))  
    .reduceByKey(lambda x, y: x + y)  
    .filter(lambda (type, count): count > 10)
```



# RDD: Fault Tolerance

RDDs track *lineage* info to rebuild lost data

```
file.map(lambda rec: (rec.type, 1))  
    .reduceByKey(lambda x, y: x + y)  
    .filter(lambda (type, count): count > 10)
```



# RDD: User Control

## **Persistence and Partitioning Strategies :**

- Indicate which RDDs they will reuse
- Choose a storage strategy for RDD  
(e.g., hint to keep in-memory storage)
- Request for RDD to be partitioned across machines (i.e., placement optimizations)

# RDD: Advantage over MapReduce

---

- MapReduce:
  - computational power of cluster, but not of distributed main memory
  - **hence, time consuming and slow**
- RDDs in contrast support:
  - in-memory storage
  - in-memory transfer of data

# RDD vs. Traditional Shared Memory

Aspect	RDDs	Distr. Shared Mem.
Reads	Coarse- or fine-grained	Fine-grained
Writes	Coarse-grained	Fine-grained
Consistency	Trivial (immutable)	Up to app / runtime
Fault recovery	Fine-grained and low-overhead using lineage	Requires checkpoints and program rollback
Straggler mitigation	Possible using backup tasks	Difficult
Work placement	Automatic based on data locality	Up to app (runtimes aim for transparency)
Behavior if not enough RAM	Similar to existing data flow systems	Poor performance (swapping?)

# Creating RDDs

---

- Loading from external dataset (file)
- Creating from another RDD (transformation)
- Parallelizing a centralized collection

# Creating RDDs

## 1. Loading an external dataset

- Most common method for creating RDDs
- Data can be located in any storage system like HDFS, Hbase , Cassandra etc.
- Example:

```
lines = spark.textFile("hdfs://...")
```



**Support for HDFS, HBase, Amazon S3, ...**

**RDD: #partitions = #of HDFS blocks**

# Creating RDDs

## 2. Creating an RDD from an Existing RDD

- An existing RDD can be used to create a new RDD.
- The Parent RDD remains intact and is not modified.
- The parent RDD can be used for further operations.
- Example

```
errors = lines.filter(_.startsWith("ERROR"))
```



**New RDD**



# Creating RDDs

## 3. Parallelizing a Central Collection

```
val data = Array(1, 2, 3, 4, 5, 100, 8, 7, ....)
```

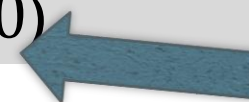
```
val distData = sc.parallelize(data)
```



**SparkContext**

```
val data = Array(1, 2, 3, 4, 5, 100, 8, 7, ....)
```

```
val distData = sc.parallelize(data, 10)
```



**Create 10  
partitions**

# Operations on RDDs

**Create new RDD**

**Return value to caller**

**No execution is triggered  
for these operations**

**Execution is triggered  
for these operations**

- **Transformation Ops.** & **Action Ops.**



Similar to map-side of Hadoop



Similar to reduce-side of Hadoop

# Transformation Ops

# Transformation Ops

---

- Operate on one RDD and generate new RDD
- The input RDD is left intact
- Lazy evaluation
- Examples: *map, filter, join*

# Transformation Ops: Example I

## Transformations

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(_.startsWith("ERROR"))
```

- Original parent RDD is left intact and can be used in future transformations.
- No action takes place, just metadata of errors RDD are created.

# Transformation Ops: Example II

```
val lines = sc.textFile("data.txt")  
val lineLengths = lines.map(s => s.length)
```



**Up to Spark to keep it in memory  
OR re-compute when needed**

```
lineLengths.persist()
```



**Ask Spark to keep this RDD in memory**

```
lineLengths.unpersist()
```



**Ask Spark to remove from cache**

---

# Action Ops

# Action Ops

- Perform a computation on existing RDDs producing a result.
- Result is either:
  - Returned to the Driver Program.
  - Stored in a files system (like HDFS).
- Examples:
  - *count()*
  - *collect()*
  - *reduce()*
  - *save()*



# Action Ops: Example I

```
val lines = sc.textFile("data.txt")
val lineLengths = lines.map(s => s.length)
val totalLength = lineLengths.reduce((a, b) => a + b)
```

Note: You can apply Reduce op on any type that has + fct.

# Action Ops: Example II [[Link](#)]

```
val logFile = "hdfs://master.backtobazics.com:9000/user/root/sample.txt"
val lineRDD = sc.textFile(logFile)
//Transformation 1 -> DAG created
//{DAG: Start -> [sc.textFile(logFile)]}

val wordRDD = lineRDD.flatMap(_.split(" "))
//Transformation 2 -> wordRDD DAG updated
//{DAG: Start -> [sc.textFile(logFile)]
//      -> [lineRDD.flatMap(_.split(" "))]}

val filteredWordRDD = wordRDD.filter(_.equalsIgnoreCase("the"))
//Transformation 3 -> filteredWordRDD DAG updated
//{DAG: Start -> [sc.textFile(logFile)]
//      -> [lineRDD.flatMap(_.split(" "))]
//      -> [wordRDD.filter(_.equalsIgnoreCase("the"))]}

filteredWordRDD.collect
//Action: collect
//Execute DAG & collect result to driver node
```

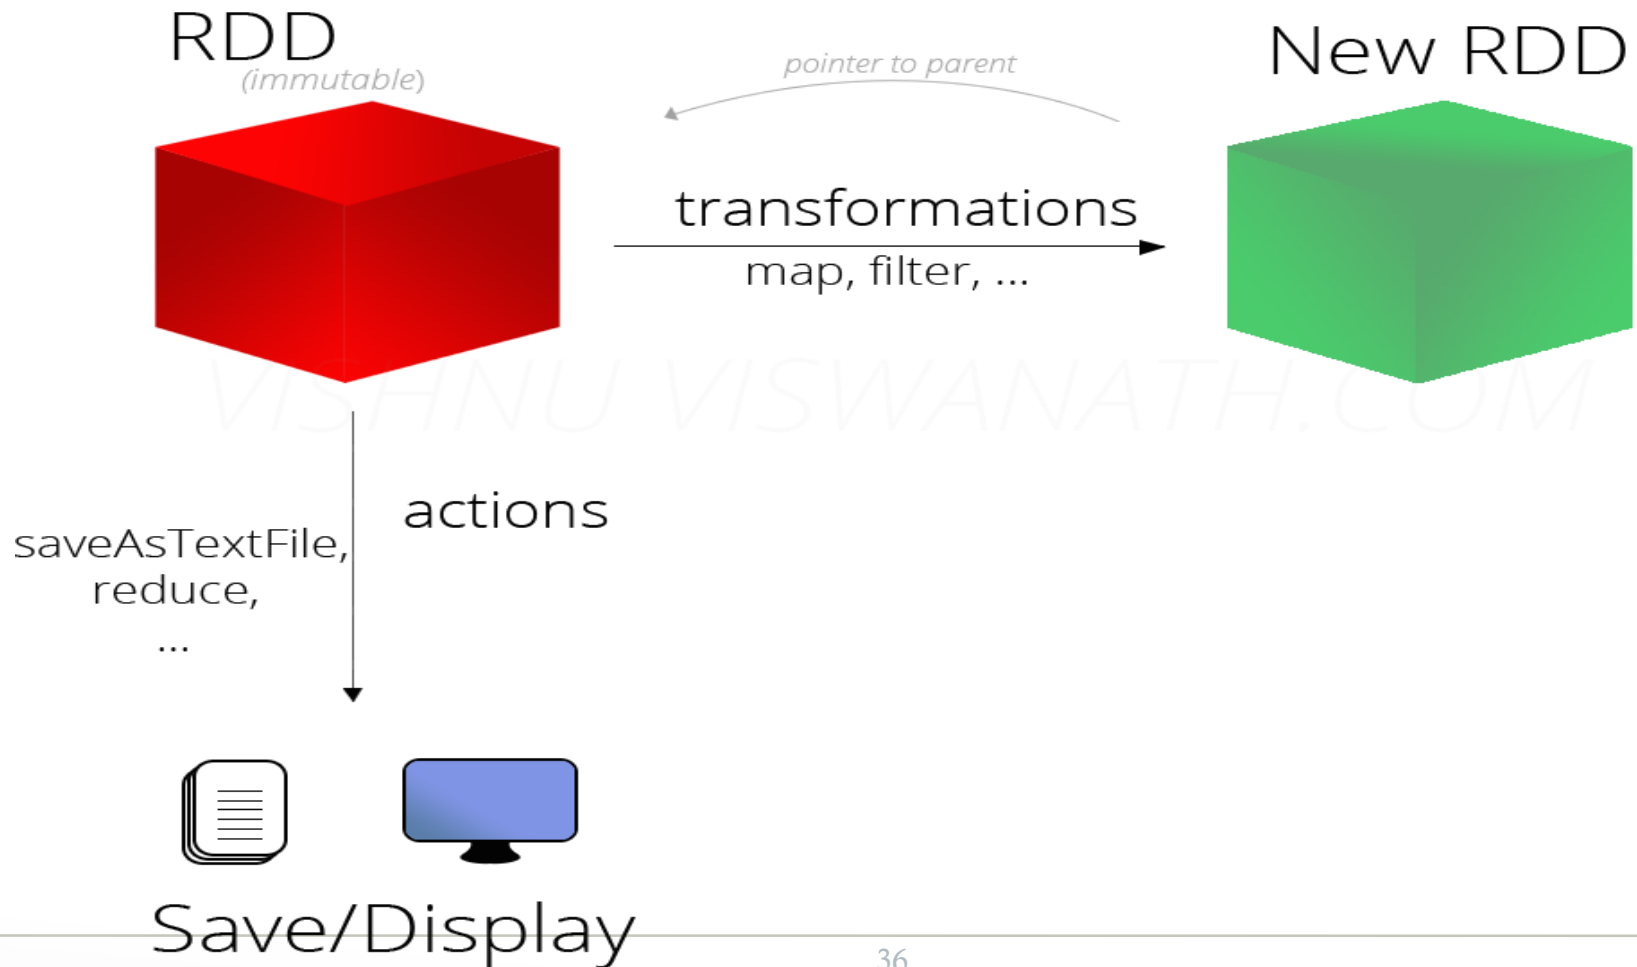
```
val logFile = "hdfs://master.backtobazics.com:9000/user/root/sample.txt"
val lineRDD = sc.textFile(logFile)
//Transformation 1 -> DAG created
//{DAG: Start -> [sc.textFile(logFile)]}

val wordRDD = lineRDD.flatMap(_.split(" "))
//Transformation 2 -> wordRDD DAG updated
//{DAG: Start -> [sc.textFile(logFile)]
//      -> [lineRDD.flatMap(_.split(" "))]}

val filteredWordRDD = wordRDD.filter(_.equalsIgnoreCase("the"))
//Transformation 3 -> filteredWordRDD DAG updated
//{DAG: Start -> [sc.textFile(logFile)]
//      -> [lineRDD.flatMap(_.split(" "))]
//      -> [wordRDD.filter(_.equalsIgnoreCase("the"))]}

filteredWordRDD.collect
//Action: collect
//Execute DAG & collect result to driver node
```

# Transformations vs. Actions



# Lazy Evaluation

- Transformation ops follow lazy evaluation
- Results not physically computed right away
- Metadata regarding transformations recorded
- Transformations are implemented only when an action is invoked

# Example: Lazy Evaluation

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(_.startsWith("ERROR"))  
errors.count()
```



**Execution is triggered here**

# RDD Fault Tolerance

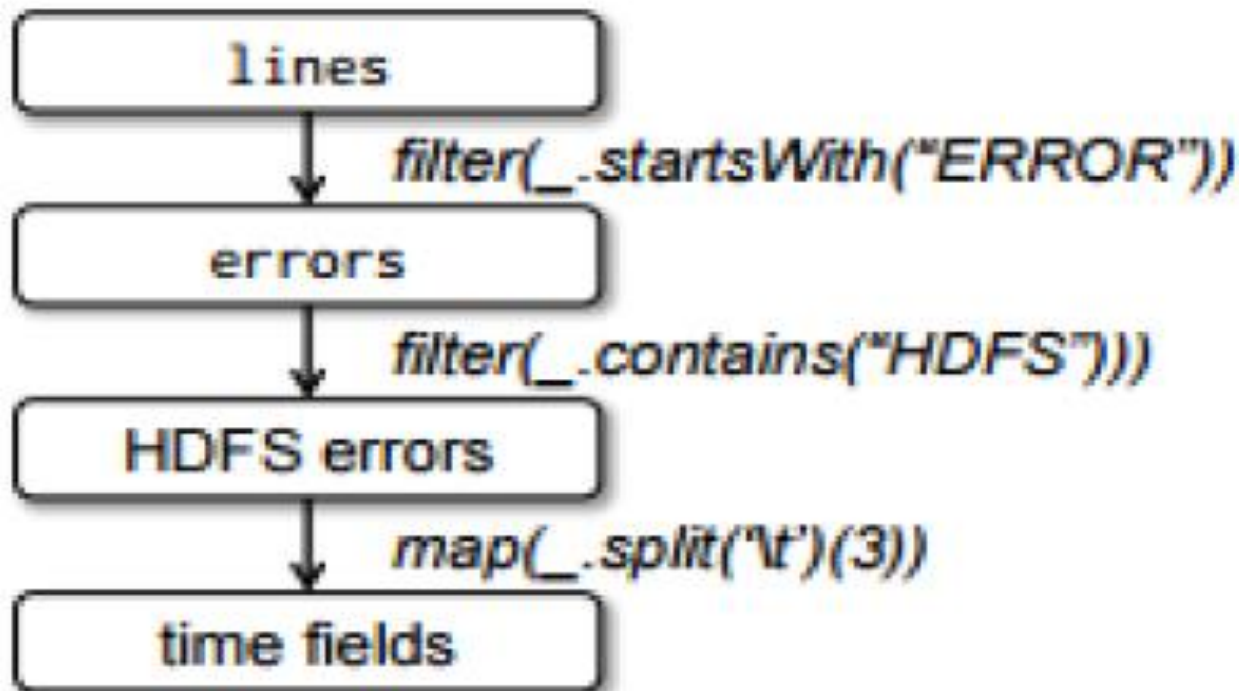
# RDD Fault Tolerance

- In-memory RDDs are not replicated
  - RAM is still limited in size (Scarce Resource)
- **Lineage Graph (“Logical Execution Plan”)**
  - Directed Acyclic Graph (DAG) produced when Sparkcontext requested to run spark job
  - Maintain dependencies between RDDs
  - Go back to closest disk-based RDD

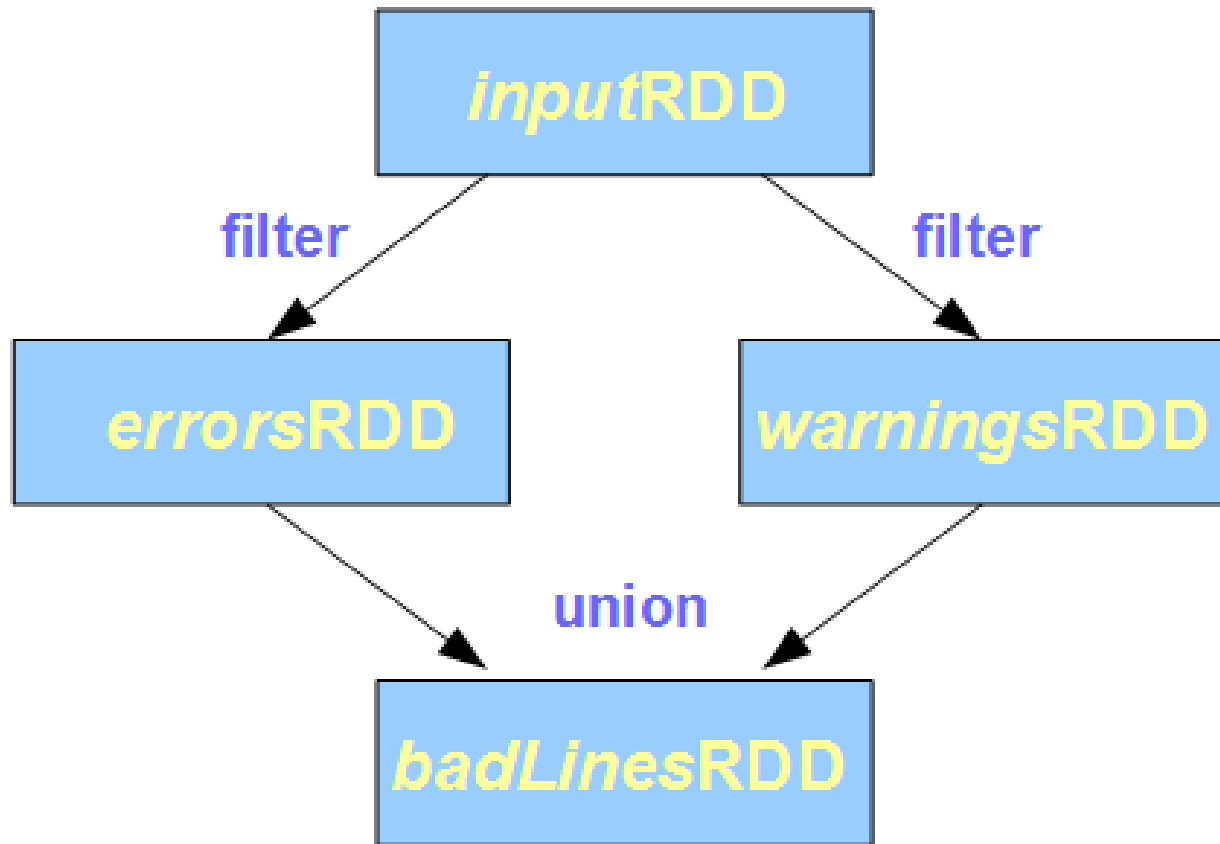


# Lineage Graph

- Not storing the data, but instead stores how it is generated (the processing steps)



# Lineage Graph



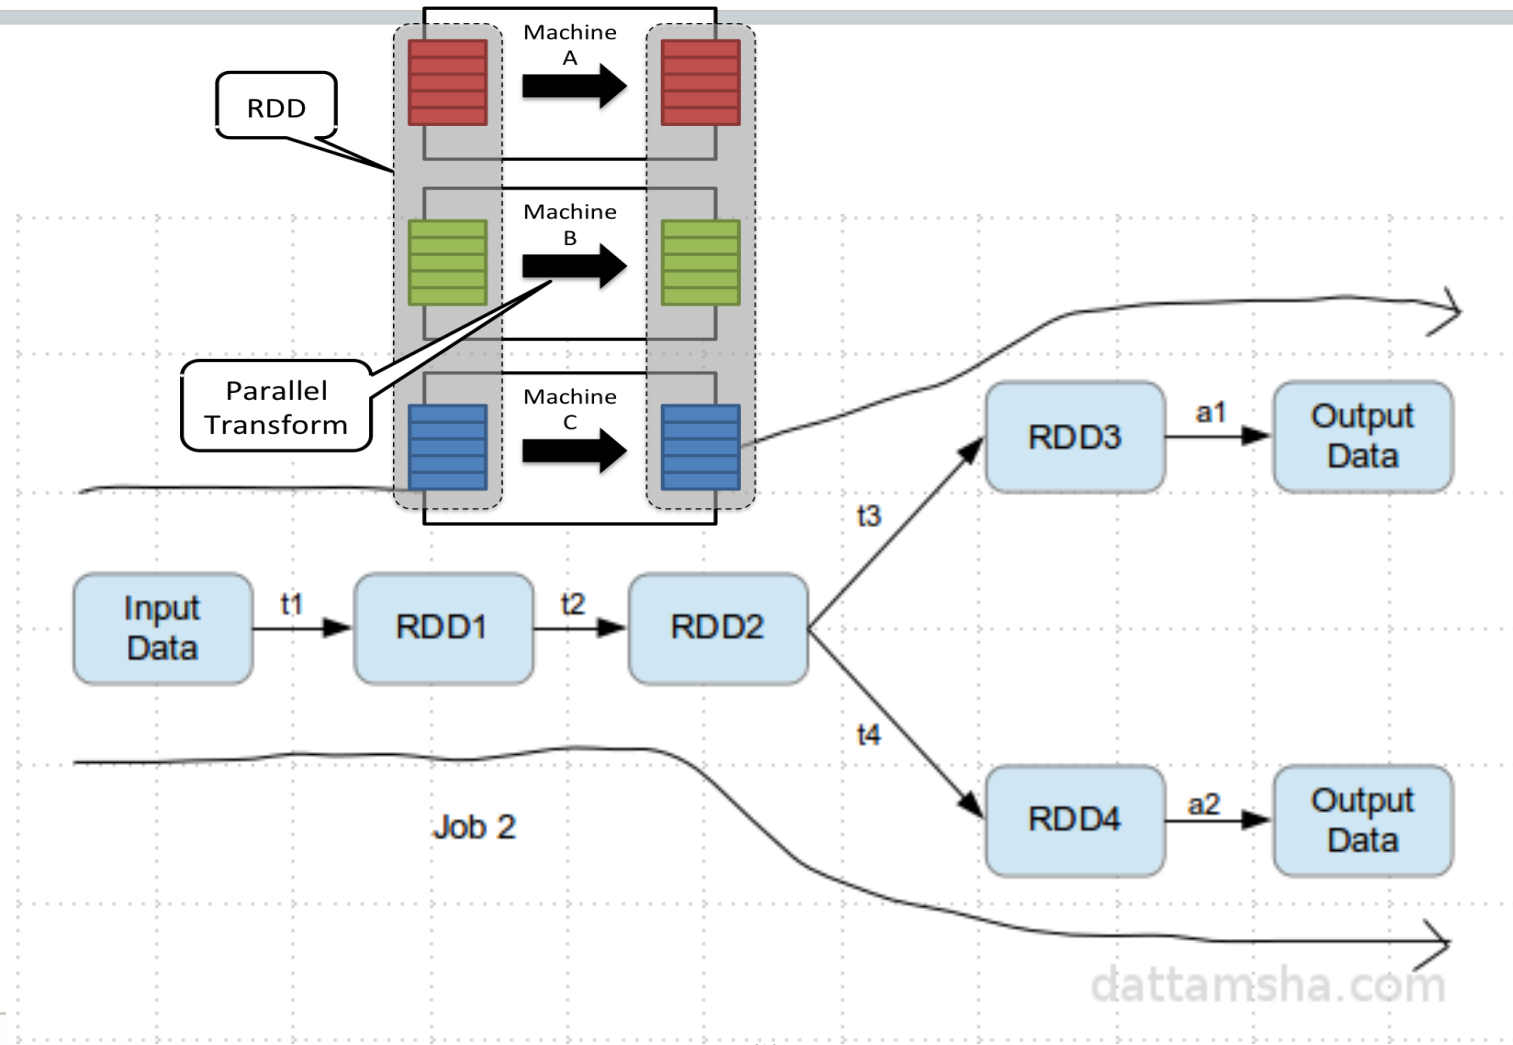
**Figure 3-1.** RDD lineage graph created during log analysis

# Representation of RDDs

# Representation of RDDs

- **Each RDD is divided into:**
  - Multiple partitions
  - Dependencies on parent RDD(s)
- **Two types of dependencies**
  - Narrow (1 to 1)
  - Wide ( many to many )

# Representation of RDDs

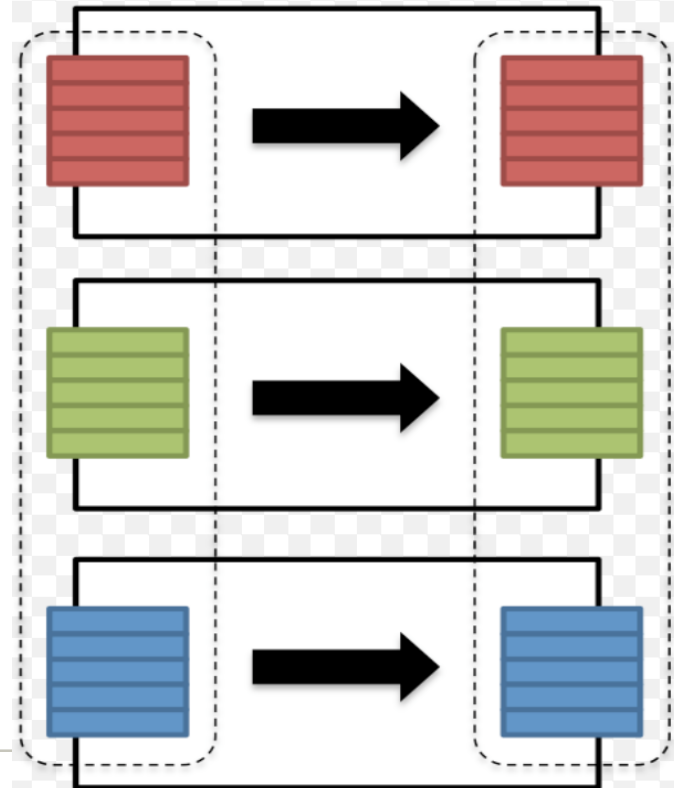


# Narrow Dependency

- 1-1 relationship between child-parent partitions
- Example Ops: *Filter* & *Map*
- Relatively cheap process

## Narrow transformation

- Input and output stays in same partition
- No data movement is needed

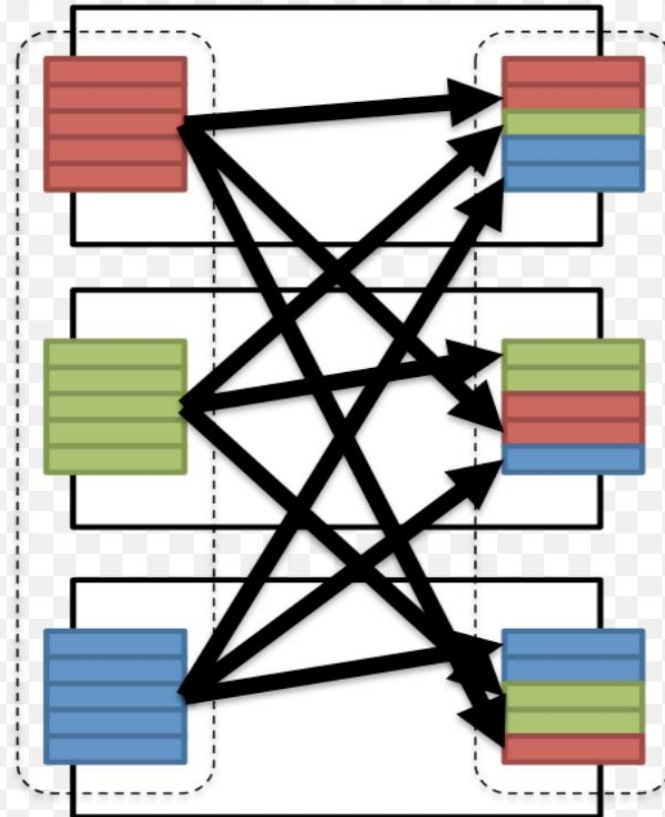


# Wide Dependency

- M-1 or M-M relationship between child-parent partitions
- Example Ops:  
*Join & Grouping*
- More expensive

## Wide transformation

- Input from other partitions are required
- Data shuffling is needed before processing



# Advanced Interfaces on RDDs

Operation	Meaning
<code>partitions()</code>	Return a list of Partition objects
<code>preferredLocations(<i>p</i>)</code>	List nodes where partition <i>p</i> can be accessed faster due to data locality
<code>dependencies()</code>	Return a list of dependencies
<code>iterator(<i>p</i>, <i>parentIters</i>)</code>	Compute the elements of partition <i>p</i> given iterators for its parent partitions
<code>partitioner()</code>	Return metadata specifying whether the RDD is hash/range partitioned

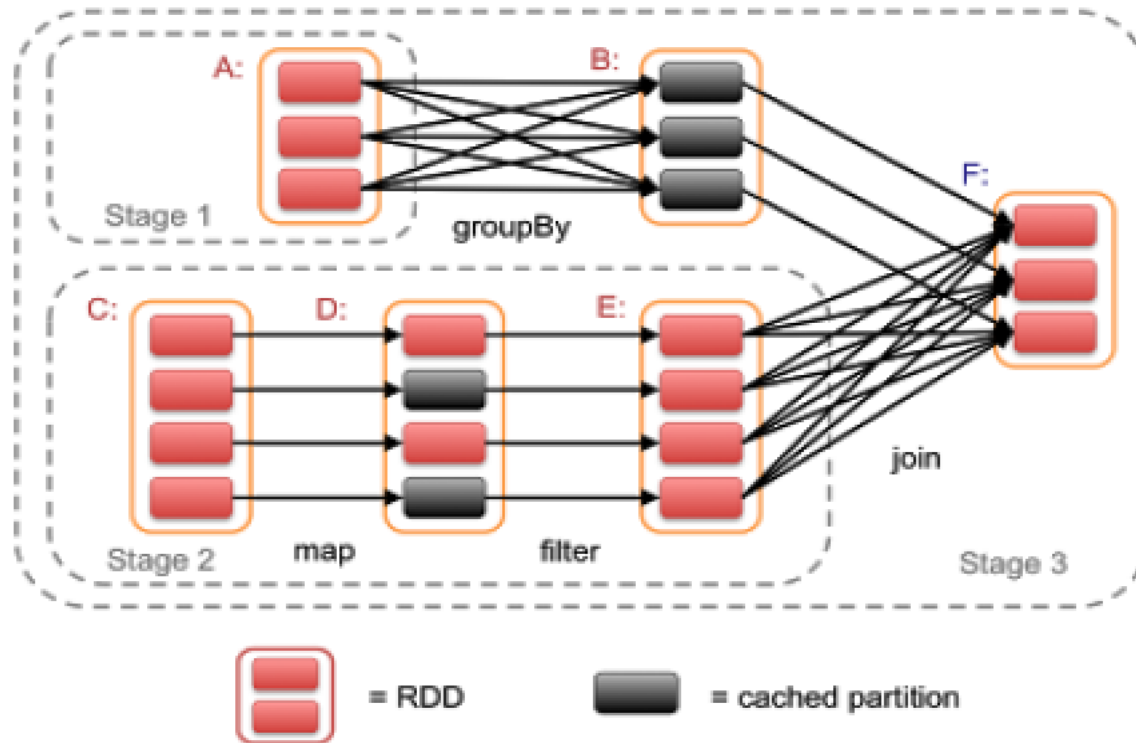
Interface used to represent RDD in Spark



# Scheduling & Memory Management

# Scheduling

- Execution is triggered when an “*Action*” op is invoke
- Scheduler checks the lineage graph to execute



# Spark Memory Management

Memory utilization is essential in Spark (Caching):

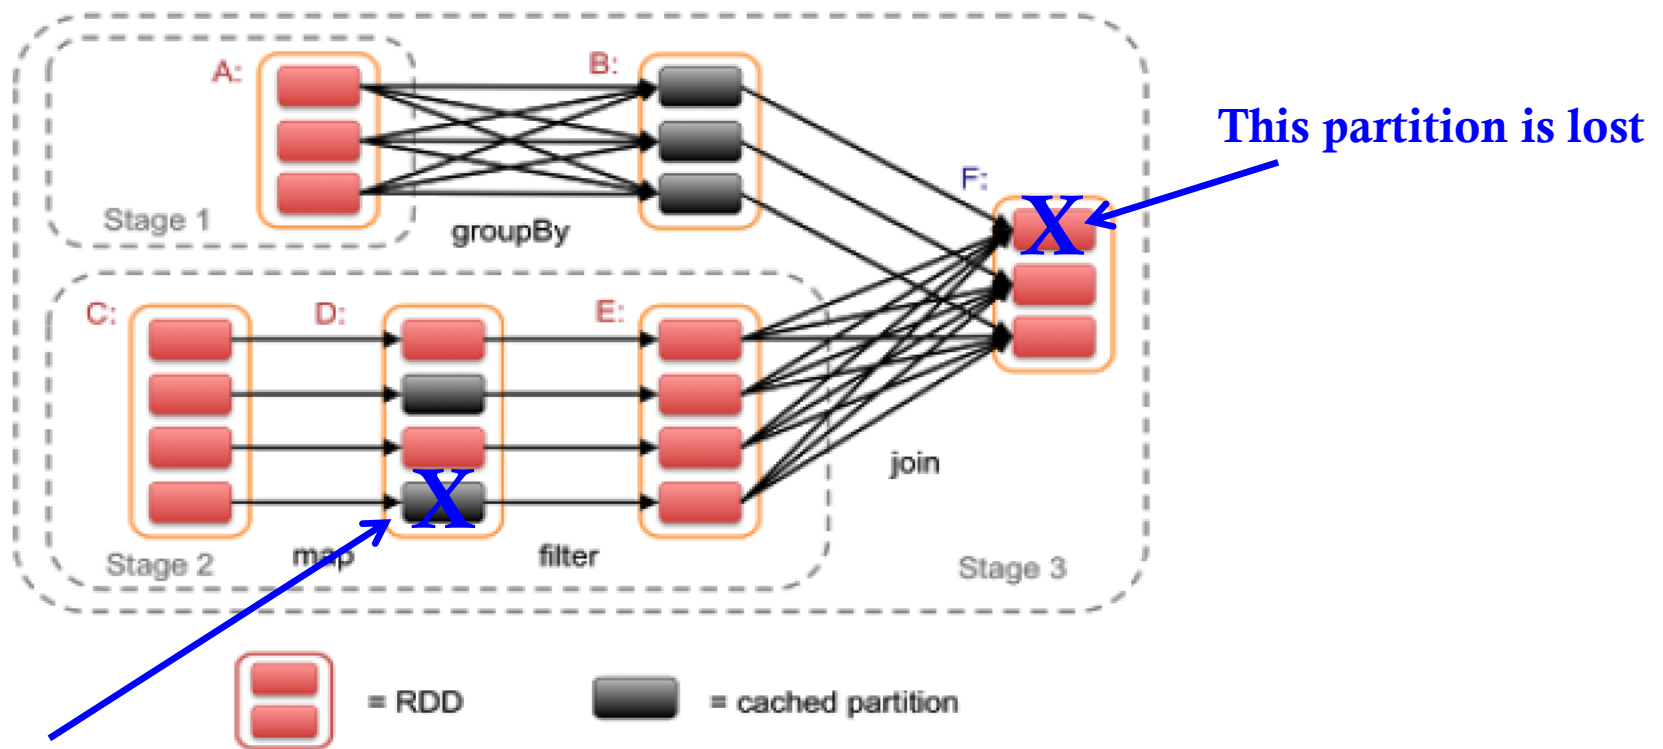
- 1. in-memory storage as (deserialized) java objects
  - - fast
- 2. in-memory storage as serialized data
  - - memory efficient, but not as fast
- 3. on-disk storage (if RDD too large to fit in RAM)
  - - slower and costly

# Replacement Policy

- LRU eviction policy at level of RDD partitions is used
- **When a new RDD partition is created**
  - If there is space in memory → Cache it
  - If not → evict one or more partitions from LRU RDD
- Use “*persistence priority*” to prevent eviction of important RDDs

# RDD Recovery

- In case of failure and losing an RDD partition



This partition is lost

# RDD Recovery

- Recovery can be time consuming for RDDs **with long lineage chains**
- Use *Checkpoint Mechanism* to make some RDDs persistent:
  - User defined, OR
  - System controlled, OR
  - More intelligent ways, e.g., workload driven

# RDD Summary

---

- Memory Cache RDD is **fault-tolerant**:
  - using sophisticated fine-grained Lineage Mechanism
- Disk Copy RDD is **fault-tolerant**:
  - using HDFS replication
- Aware User: Can manually call **unpersist()**

---

# Summary



# Example I

## Example: Mining Console Logs

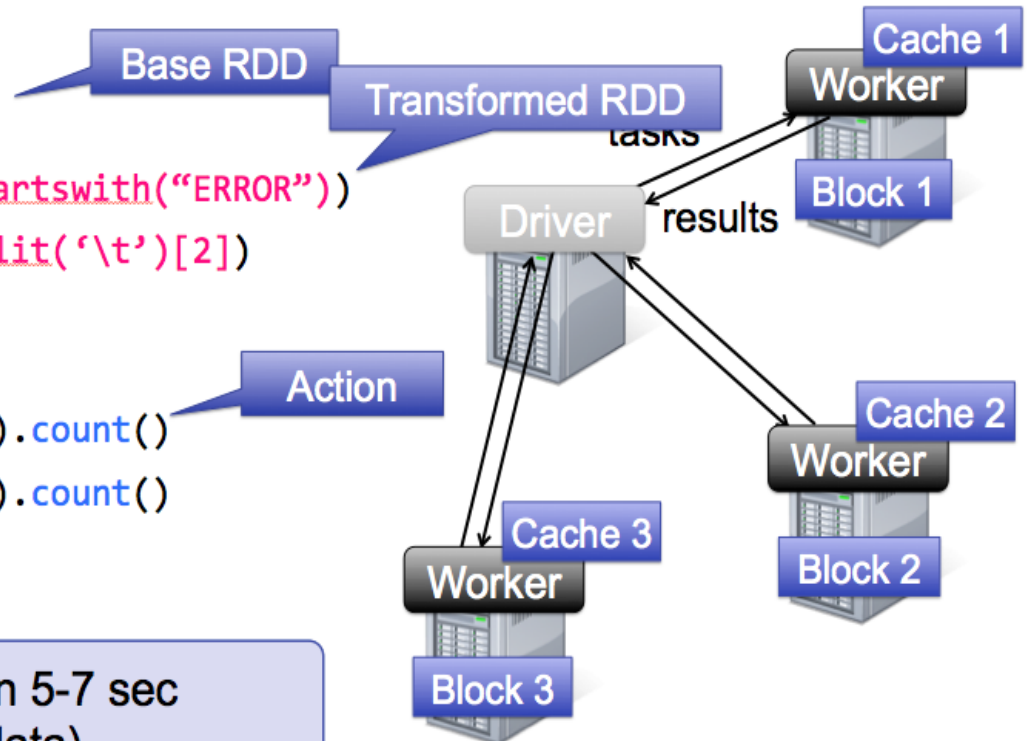
- Load error messages from a log into memory, then interactively search for patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split('\t')[2])
messages.cache()
```

```
messages.filter(lambda s: "foo" in s).count()
messages.filter(lambda s: "bar" in s).count()
```

...

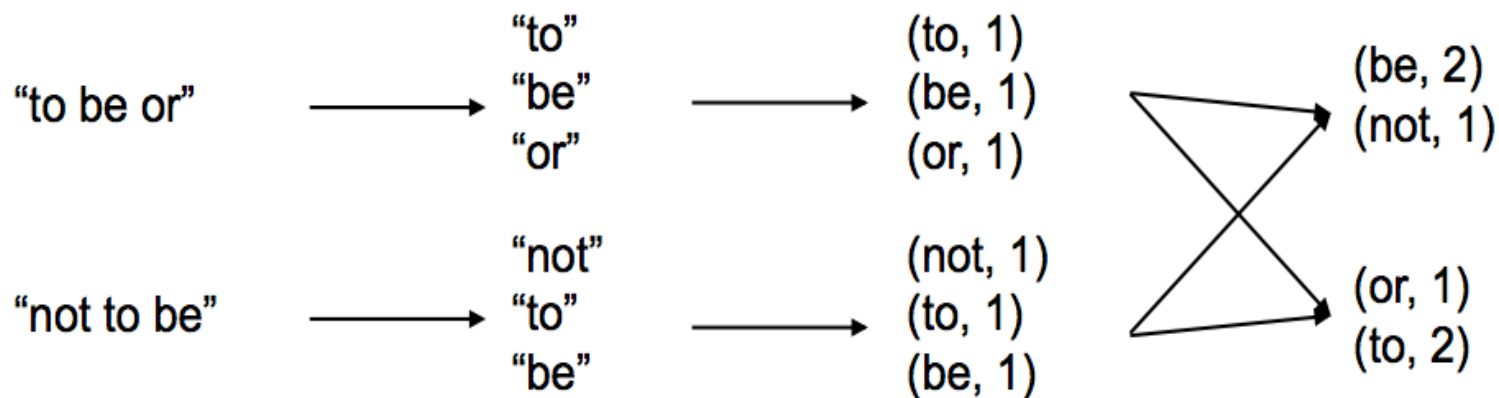
**Result:** scaled to 1 TB data in 5-7 sec  
(vs 170 sec for on-disk data)



# Example II

## Example: Word Count

```
lines = sc.textFile("hamlet.txt")
counts = lines.flatMap(lambda line: line.split(" ")) \
               .map(lambda word: (word, 1)) \
               .reduceByKey(lambda x, y: x + y)
```

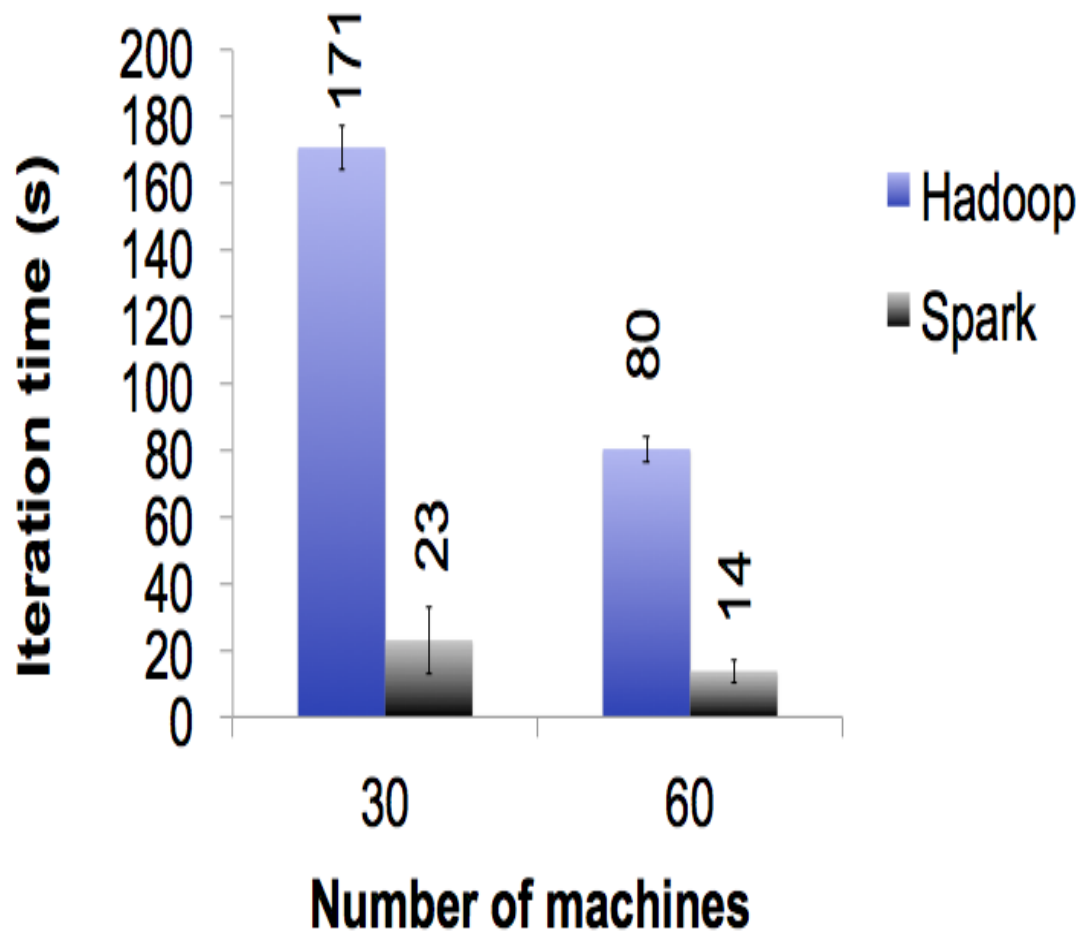


# Performance

---

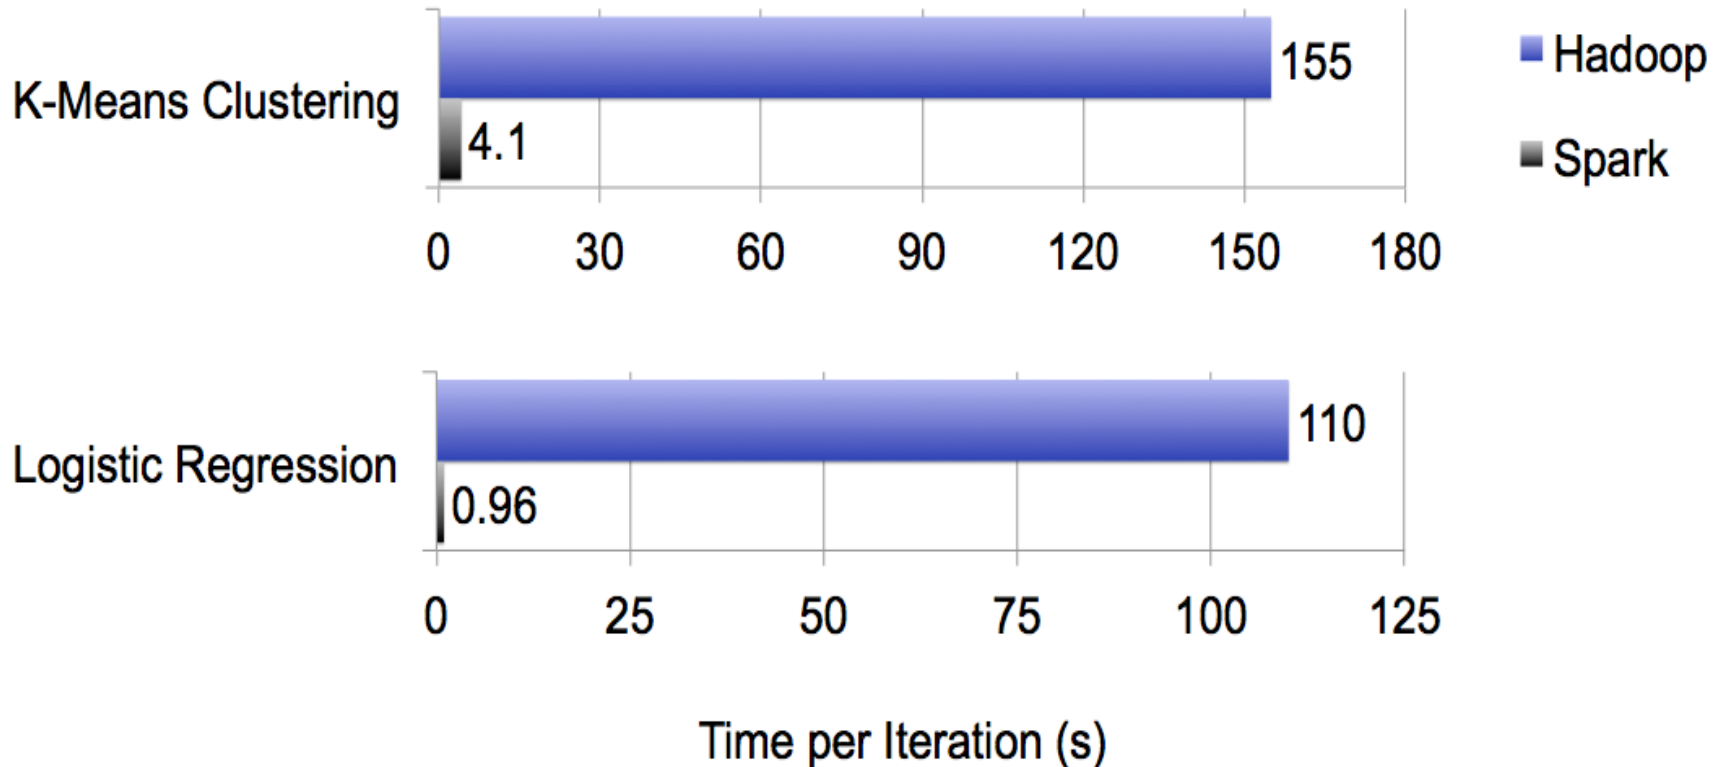
- Spark is faster than Hadoop 10x – 100x
- Especially for iterative algorithms

# PageRank Performance



# K-Means & L. Regression

## Other Iterative Algorithms



# Performance (Revisited)

## POSITIVE:

- Spark is faster than Hadoop 10x – 100x
- Especially for iterative algorithms

## NEGATIVE:

- Spark needs much more memory than Hadoop
- Spark performance is more sensitive to availability of memory

## SAME:

- *Jobs that require one iteration:*
  - No big difference between Spark & Hadoop